

In [1]:

```
import numpy as np
import torch
import matplotlib.pyplot as plt
import torchvision
import torchvision.transforms as transforms
from torch import nn
```

In [20]:

```
#HYPERPARAMETERS
learning_rate=0.001
epochs=10
batch_size= 300
number_of_pc=30 #number of principal components (for PCA)
```

In [3]:

```
train_data = torchvision.datasets.MNIST(root="./",train=True,transform=transform
s.ToTensor(),download=True)
test_data = torchvision.datasets.MNIST(root="./",train=False,transform=transform
s.ToTensor(),download=True)
```

In [4]:

```
train_loader = torch.utils.data.DataLoader(dataset=train_data,batch_size=batch_s
ize,shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_data,batch_size=len(test_
data),shuffle=False)
```

In [5]:

```
test_sample_loader = torch.utils.data.DataLoader(dataset=test_data.data[9705:971
5],shuffle=False,batch_size=10)
```

In [6]:

```
train_dataset=train_data.data.reshape(train_data.data.shape[0],train_data.data.s
hape[1]*train_data.data.shape[2])
test_dataset = test_data.data.reshape(test_data.data.shape[0],test_data.data.sha
pe[1]*test_data.data.shape[2])
test_dataset_sampled= test_dataset[np.arange(9705,9715),:]
```

Question 1

In [9]:

```
def PCA(input_data,top_k_ev):
    '''
    function: This function performs pca.
    Input: input_data = (torch matrix) = of shape num_datapts,784pixels
           top_k_ev = (integer) = how many principal components to be taken

    Output: top_k_eigen_vectors= shape(784,top_k_ev) top k eigen vectors in column
    s
           centered_ip_data= shape(num_datapts,784) = centered ip data (used for
    reconstruction from principal components)

    '''

    input_mean = torch.mean(input_data,0)
    centered_ip_data = input_data-input_mean
    cov_matrix = torch.matmul(centered_ip_data.T,centered_ip_data)
    eigen_values, eigen_vectors = torch.linalg.eigh(cov_matrix)
    eigen_values_descending,indices = torch.sort(eigen_values,descending=True)
    top_k_eigen_values,top_k_indices = eigen_values_descending[:top_k_ev],indices[:top_k_ev]
    top_k_eigen_vectors = eigen_vectors[:,top_k_indices]

    assert top_k_eigen_vectors.shape == (784,top_k_ev)
    assert centered_ip_data.shape == input_data.shape
    assert input_mean.shape == torch.Size([784])

    return top_k_eigen_vectors
```

In [10]:

```
pc=PCA(train_dataset.float(),number_of_pc)
```

In [11]:

```
def reconstruct_data(principal_components,dataset):
    '''
    function: This function reconstructs the datapoints in lower dimension (i.e. top k eigen_vector components).

    Input: principal_components= shape(784,top_k_ev) top k eigen vectors in column
    s
           centered_ip_data= shape(num_datapts,784) = centered ip data (used for
    reconstruction from principal components)

    Output: projected_data = (torch matrix) = of shape num_datapts , top_k_ev

    '''
    projection_matrix = torch.matmul(principal_components,principal_components.T)
    projected_data = torch.matmul(dataset,projection_matrix)

    assert projected_data.shape == (dataset.shape[0],principal_components.shape[0])

    return projected_data
```

In [12]:

```
reconstructed_test_data_sampled=reconstruct_data(pc,test_dataset_sampled.float
())
```

In [13]:

```
class AE1(nn.Module):
    def __init__(self):
        super(AE1, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784,512),
            nn.ReLU(),
            nn.Linear(512,256),
            nn.ReLU(),
            nn.Linear(256,128),
            nn.ReLU(),
            nn.Linear(128,30),
            nn.ReLU())
        self.decoder =nn.Sequential(
            nn.Linear(30,128),
            nn.ReLU(),
            nn.Linear(128,256),
            nn.ReLU(),
            nn.Linear(256,784),
            nn.ReLU())

    def forward(self,x):
        x=self.encoder(x)
        encoded_output=x
        x=self.decoder(x)
        return x,encoded_output
```

In [14]:

```

model1=AE1()
criterion1=nn.MSELoss()
optimizer1 = torch.optim.Adam(model1.parameters(),lr=learning_rate)

training_loss=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,_=model1(images)
        loss=criterion1(outputs,images)
        training_loss.append(loss.item())

        optimizer1.zero_grad()
        loss.backward()
        optimizer1.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")

```

```

Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed


```

In [15]:

```

plt.figure(figsize=(12,6))
plt.plot(range(1,len(training_loss)+1),training_loss)
plt.title("Training Loss of AE1 model vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.grid()

```



In [16]:

```

#MODEL EVALUATION AND RESULT PLOTTING
modell.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs,_ = modell(images.float())

plt.rcParams["figure.figsize"] = (12,6)
for i in range (10):
    fig, (ax1, ax2, ax3) = plt.subplots(1,3)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('AE Reconstructed Image')
    ax2.axis("off")


    ax3.imshow(reconstructed_test_data_sampled[i].reshape(28,28),cmap='gray')
    ax3.set_title('PCA Reconstructed Image')
    ax3.axis("off")
    print("_____")
    print("Reconstruction Error in AE:",np.dot(((images[i].detach().numpy()/255)-(
outputs[i].detach().numpy()/255)),((images[i].detach().numpy()/255)-(outputs[i].
detach().numpy()/255)).T))
    print("Reconstruction Error in PCA:",np.dot(((images[i].detach().numpy()/255)-(
reconstructed_test_data_sampled[i].detach().numpy()/255)),((images[i].detach().
numpy()/255)-(reconstructed_test_data_sampled[i].detach().numpy()/255)).T))
    plt.show()

```


Reconstruction Error in AE: 2.7472248430821784

Reconstruction Error in PCA: 4.90691098406753


Original Image



AE Reconstructed Image



PCA Reconstructed Image




Reconstruction Error in AE: 21.762326668644867

Reconstruction Error in PCA: 16.86237506274601

Original Image



AE Reconstructed Image



PCA Reconstructed Image



Reconstruction Error in AE: 21.333864113817814

Reconstruction Error in PCA: 16.045755599921307


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 14.849075947558745

Reconstruction Error in PCA: 10.714795712809165


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 15.122479463273889
Reconstruction Error in PCA: 14.848193698340676


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 22.65667615807221
Reconstruction Error in PCA: 20.390818940836983


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 16.34095082586718
Reconstruction Error in PCA: 11.735863192062858


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 15.306385477625026

Reconstruction Error in PCA: 14.35657074432876


Original Image



AE Reconstructed Image




PCA Reconstructed Image




Reconstruction Error in AE: 14.786297900784824

Reconstruction Error in PCA: 13.215489962853521

Original Image



AE Reconstructed Image




PCA Reconstructed Image



Reconstruction Error in AE: 19.088079263514267

Reconstruction Error in PCA: 16.605796880360153



Observations:

- Visually, Reconstructed Autoencoder images look appealing as they have better contrast.
- Reconstruction error-wise, Reconstructed PCA images are found to have lesser error (doesnot have much difference though), however they lack the contrast.

Question 2

In [21]:

```
class AE2(nn.Module):
    def __init__(self, hid):
        super(AE2, self).__init__()
        self.hid=hid
        self.encoder = nn.Sequential(
            nn.Linear(784, self.hid),
            nn.ReLU()
        )
        self.decoder =nn.Sequential(
            nn.Linear(self.hid, 784),
            nn.ReLU()
        )

    def forward(self, x):
        x=self.encoder(x)
        encoded_output=x
        x=self.decoder(x)
        return x, encoded_output
```

In [18]:

```
#HIDDEN_SIZE=64
model_hid64=AE2(64)
criterion_hid64=nn.MSELoss()
optimizer_hid64 = torch.optim.Adam(model_hid64.parameters(),lr=learning_rate)

training_loss_hid64=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,_=model_hid64(images)
        loss=criterion_hid64(outputs,images)
        training_loss_hid64.append(loss.item())

        optimizer_hid64.zero_grad()
        loss.backward()
        optimizer_hid64.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [19]:

```
#HIDDEN_SIZE=128
model_hid128=AE2(128)
criterion_hid128=nn.MSELoss()
optimizer_hid128 = torch.optim.Adam(model_hid128.parameters(),lr=learning_rate)

training_loss_hid128=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,_=model_hid128(images)
        loss=criterion_hid128(outputs,images)
        training_loss_hid128.append(loss.item())

        optimizer_hid128.zero_grad()
        loss.backward()
        optimizer_hid128.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [22]:

```
#HIDDEN_SIZE=256
model_hid256=AE2(256)
criterion_hid256=nn.MSELoss()
optimizer_hid256 = torch.optim.Adam(model_hid256.parameters(),lr=learning_rate)

training_loss_hid256=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,_=model_hid256(images)
        loss=criterion_hid256(outputs,images)
        training_loss_hid256.append(loss.item())


        optimizer_hid256.zero_grad()
        loss.backward()
        optimizer_hid256.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [21]:

```
plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid64,label="Hidden S  
ize=64")  
plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid128,label="Hidden  
Size=128")  
plt.plot(range(1,len(training_loss_hid64)+1),training_loss_hid256,label="Hidden  
Size=256")  
plt.legend()  
plt.grid()  
plt.title("Training Loss Vs Iterations")  
plt.xlabel("Number of Iterations")  
plt.ylabel("Training Loss")  
plt.show()
```



In [22]:

```

model_hid64.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid64,_ = model_hid64(images.float())

model_hid128.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid128,_ = model_hid128(images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        outputs_hid256,activations_hid256 = model_hid256(images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
if i==5:
    fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_hid64[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('AE_hid64 Reconstructed Image')
    ax2.axis("off")

    ax3.imshow(outputs_hid128[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('AE_hid128 Reconstructed Image')
    ax3.axis("off")

    ax4.imshow(outputs_hid256[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('AE_hid256 Reconstructed Image')
    ax4.axis("off")

    print("Reconstruction Error in AE_hid64:",np.dot(((images[i].detach().numpy()/
255)-(outputs_hid64[i].detach().numpy()/255)),((images[i].detach().numpy()/255)-
(outputs_hid64[i].detach().numpy()/255)).T))
    print("Reconstruction Error in AE_hid128:",np.dot(((images[i].detach().numpy()/
/255)-(outputs_hid128[i].detach().numpy()/255)),((images[i].detach().numpy()/255
)-(outputs_hid128[i].detach().numpy()/255)).T))
    print("Reconstruction Error in AE_hid256:",np.dot(((images[i].detach().numpy()/
/255)-(outputs_hid256[i].detach().numpy()/255)),((images[i].detach().numpy()/255
)-(outputs_hid256[i].detach().numpy()/255)).T))


```

Reconstruction Error in AE_hid64: 17.168603545148102
Reconstruction Error in AE_hid128: 10.844266659450152
Reconstruction Error in AE_hid256: 9.580246903034212


Original Image




AE_hid64 Reconstructed Image



AE_hid128 Reconstructed Image



AE_hid256 Reconstructed Image



In [23]:

```

#FASHION MNIST OUTPUT
test_data_fashion = torchvision.datasets.FashionMNIST(root=".",train=False,trans
form=transforms.ToTensor(),download=True)
fashion_image_sample=test_data_fashion.data[10]

model_hid64.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1,28*28)
    outputs_hid64,_ = model_hid64(images.float())

model_hid128.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1,28*28)
    outputs_hid128,_ = model_hid128(images.float())

model_hid256.eval()
with torch.no_grad():
    images = fashion_image_sample.reshape(1,28*28)
    outputs_hid256,_ = model_hid256(images.float())

plt.rcParams["figure.figsize"] = (15,6)

fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)


ax1.imshow(fashion_image_sample.detach().numpy().reshape(28,28),cmap='gray')
ax1.set_title('Original Random Image')
ax1.axis("off")

ax2.imshow(outputs_hid64.detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('AE_hid64 Reconstructed Image')
ax2.axis("off")

ax3.imshow(outputs_hid128.detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('AE_hid128 Reconstructed Image')
ax3.axis("off")

ax4.imshow(outputs_hid256.detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('AE_hid256 Reconstructed Image')
ax4.axis("off")
plt.show()

```



In [24]:

```

#OUTPUTS FOR RANDOM IMAGE
torch.manual_seed(0)
random_image=torch.randint(low=0, high=255,size=(1,28,28))

model_hid64.eval()
with torch.no_grad():
    images = random_image.reshape(1,28*28)
    outputs_hid64,_ = model_hid64(images.float())

model_hid128.eval()
with torch.no_grad():
    images = random_image.reshape(1,28*28)
    outputs_hid128,_ = model_hid128(images.float())

model_hid256.eval()
with torch.no_grad():
    images = random_image.reshape(1,28*28)
    outputs_hid256,_ = model_hid256(images.float())

plt.rcParams["figure.figsize"] = (15,6)

fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)


ax1.imshow(random_image.numpy().reshape(28,28),cmap='gray')
ax1.set_title('Original Random Image')
ax1.axis("off")

ax2.imshow(outputs_hid64.detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('AE_hid64 Reconstructed Image')
ax2.axis("off")

ax3.imshow(outputs_hid128.detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('AE_hid128 Reconstructed Image')
ax3.axis("off")

ax4.imshow(outputs_hid256.detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('AE_hid256 Reconstructed Image')
ax4.axis("off")
plt.show()

```



Observations:

- For an image sample from MNIST dataset, AE with hidden size 256 seems to reconstruct the image in a better manner. This was expected as more components would be able to reproduce the image in a better manner by capturing the minor details of the input image. This is also proved by the lesser reconstruction error by this model.
- For an image from Fashion MNIST dataset, the models try to reconstruct the image but the models with 64 and 128 hidden size could not reconstruct properly. On the other hand, 256 hidden size model does better job than them but also seems to struggle in the reconstruction, though if the model was trained on fashion mnist dataset, it would've performed better.
- For a random image, the models perform as if they were trying to find out the digits, so the center pixels have some noisy output and the outer/edge pixels seem to be off (similar to the MNIST digit dataset, as our model has been trained on them.)

Question 3

In [25]:

```
class AE3_SparseAutoencoder(nn.Module):
    def __init__(self):
        super(AE3_SparseAutoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784, 1156),
            nn.ReLU())
        self.decoder = nn.Sequential(
            nn.Linear(1156, 784),
            nn.ReLU())

    def forward(self, x):
        x = self.encoder(x)
        encoded_output = x
        l1_norm = torch.norm(x, p=1)
        x = self.decoder(x)
        return x, l1_norm, encoded_output
```

In [26]:

```
lambda_ = 3*1e-6
model_3_a=AE3_SparseAutoencoder()
criterion_3_a=nn.MSELoss()
optimizer_3_a = torch.optim.Adam(model_3_a.parameters(),lr=learning_rate)

training_loss_3_a=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,l1_norm,_=model_3_a(images)
        loss=criterion_3_a(outputs,images)+lambda_*l1_norm
        training_loss_3_a.append(loss.item())

    optimizer_3_a.zero_grad()
    loss.backward()
    optimizer_3_a.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [27]:

```
lambda_ = 1.5*1e-6
model_3_b=AE3_SparseAutoencoder()
criterion_3_b=nn.MSELoss()
optimizer_3_b = torch.optim.Adam(model_3_b.parameters(),lr=learning_rate)

training_loss_3_b=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,l1_norm,_=model_3_b(images)
        loss=criterion_3_b(outputs,images)+lambda_*l1_norm
        training_loss_3_b.append(loss.item())

        optimizer_3_b.zero_grad()
        loss.backward()
        optimizer_3_b.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [28]:

```
lambda_ = 0.6*1e-6
model_3_c=AE3_SparseAutoencoder()
criterion_3_c=nn.MSELoss()
optimizer_3_c = torch.optim.Adam(model_3_c.parameters(),lr=learning_rate)

training_loss_3_c=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        outputs,l1_norm,_=model_3_c(images)
        loss=criterion_3_c(outputs,images)+lambda_*l1_norm
        training_loss_3_c.append(loss.item())

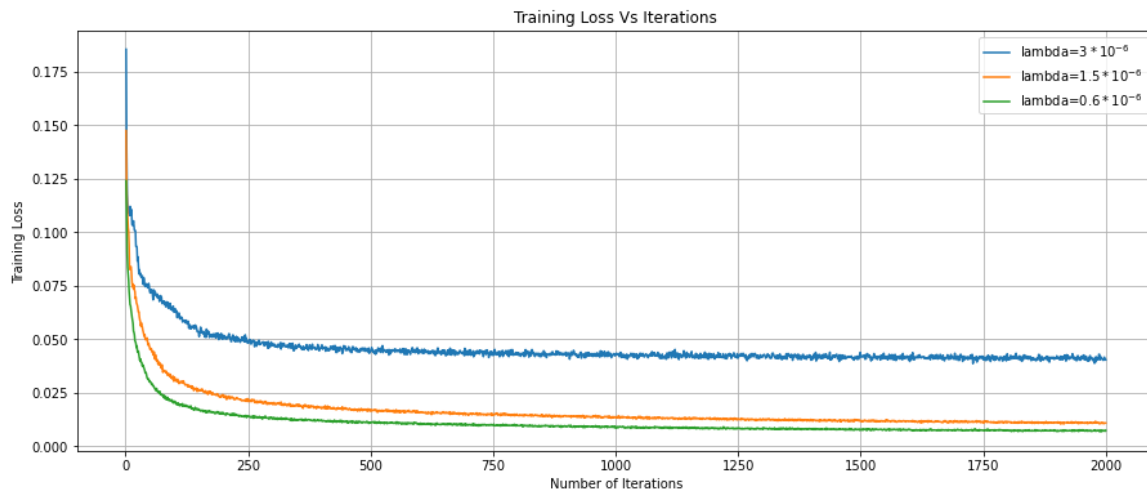
        optimizer_3_c.zero_grad()
        loss.backward()
        optimizer_3_c.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [29]:

```
plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_a,label="lambda=$3*10^{ -6}$")
plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_b,label="lambda=$1.5*10^{ -6}$")
plt.plot(range(1,len(training_loss_3_a)+1),training_loss_3_c,label="lambda=$0.6*10^{ -6}$")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()
```



In [30]:

```

model_3_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_a,_,activation_3_a = model_3_a(images.float())

model_3_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_b,_,activation_3_b = model_3_b(images.float())

model_3_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        outputs_3_c,_,activation_3_c= model_3_c(images.float())

plt.rcParams["figure.figsize"] = (15,6)

for i in range(10):
    fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_3_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Sparse AE lambda=$3*10^{-6}$')
    ax2.axis("off")


    ax3.imshow(outputs_3_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Sparse AE lambda=$1.5*10^{-6}$')
    ax3.axis("off")

    ax4.imshow(outputs_3_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Sparse AE lambda=$0.6*10^{-6}$')
    ax4.axis("off")


    plt.show()

    print("Reconstruction Error in SparseAE lambda = 3*1e-6:",np.dot(((images[i].d
detach().numpy()/255)-(outputs_3_a[i].detach().numpy()/255)),((images[i].detach()
.numpy()/255)-(outputs_3_a[i].detach().numpy()/255)).T))
    print("Reconstruction Error in SparseAE lambda = 1.5*1e-6:",np.dot(((images[i]
.detach().numpy()/255)-(outputs_3_b[i].detach().numpy()/255)),((images[i].detach
().numpy()/255)-(outputs_3_b[i].detach().numpy()/255)).T))
    print("Reconstruction Error in SparseAE lambda = 0.6*1e-6:",np.dot(((images[i]
.detach().numpy()/255)-(outputs_3_c[i].detach().numpy()/255)),((images[i].detach
().numpy()/255)-(outputs_3_c[i].detach().numpy()/255)).T))
    print("_____")


```




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 14.973617678392557
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 2.61367661274686
 03
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 3.24840965436654
 15




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 31.45378367254754
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 10.7902274525253
 57
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 10.5296989239763
 87




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 35.407050044721316
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 8.96031820622504
 3
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 13.3052458512635
 57




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 28.65200092778902
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 5.53092973497781
 6
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 5.76904392243360
 2




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 35.12167571625308
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 8.69297798002163
 3
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 7.52081489900155
 9




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 60.559072609549965
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 12.4697726663372
 5
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 14.4455125669140
 54




Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 29.60746996381112
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 6.78235910911763
 4
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 8.19118042354352



Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 34.148959339821836
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 8.82763069986666
 5
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 8.81962174996091
 4



Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 29.96832909700308
 Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 9.05938606700964
 Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 8.93139726212009



Reconstruction Error in SparseAE $\lambda = 3 \times 10^{-6}$: 30.037623122081364
Reconstruction Error in SparseAE $\lambda = 1.5 \times 10^{-6}$: 7.42948966445585
95
Reconstruction Error in SparseAE $\lambda = 0.6 \times 10^{-6}$: 11.4139420554862
81

In [31]:

```
##VISUALISING ACTIVATIONS
```

```
plt.rcParams["figure.figsize"] = (15,6)
for i in range(10):
    fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(1,5)
    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")


    ax2.imshow(np.array(activations_hid256.detach().numpy())[i].reshape(int(np.sqrt(256)),int(np.sqrt(256))), cmap='gray')
    ax2.set_title('VaniallaAE')
    ax2.axis("off")




    ax3.imshow(np.array(activation_3_a.detach().numpy())[i].reshape(int(np.sqrt(1156)),int(np.sqrt(1156))), cmap='gray')
    ax3.set_title('SparseAE lambda = $3*10^{-6}$')
    ax3.axis("off")

    ax4.imshow(np.array(activation_3_b.detach().numpy())[i].reshape(int(np.sqrt(1156)),int(np.sqrt(1156))), cmap='gray')
    ax4.set_title('SparseAE lambda=$1.5*10^{-6}$')
    ax4.axis("off")

    ax5.imshow(np.array(activation_3_c.detach().numpy())[i].reshape(int(np.sqrt(1156)),int(np.sqrt(1156))), cmap='gray')
    ax5.set_title('SparseAE lambda=$0.6*10^{-6}$')
    ax5.axis("off")

plt.show()
```





In [33]:

```
##VISUALISING ENCODER WEIGHTS
```

```
plt.rcParams["figure.figsize"] = (15,6)


fig, (ax1, ax2, ax3, ax4) = plt.subplots(1,4)
ax1.imshow(model_hid256.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
ax1.set_title('VaniallaAE')
ax1.axis("off")

ax2.imshow(model_3_a.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
ax2.set_title('SparseAE lambda =  $3 \times 10^{-6}$ ')
ax2.axis("off")

ax3.imshow(model_3_b.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
ax3.set_title('SparseAE lambda =  $1.5 \times 10^{-6}$ ')
ax3.axis("off")

ax4.imshow(model_3_c.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
ax4.set_title('SparseAE lambda =  $0.6 \times 10^{-6}$ ')
ax4.axis("off")

plt.show()
```



Observations:

- As lambda(sparsity factor) is decreased, reconstruction error goes down.
- If the lambda is continually decreased beyond a point then the model starts to overfit data. since the model is overcomplete, we need to iterate over the optimal lambda parameter to get desired output.
- as we decrease lambda(sparsity parameter) more neurons get activated. As can be observed from activation visualisations. However, in vanilla AE all the neurons are activated.
- Visualisation of encoder weights is difficult as nothing could be inferred from the images.

Question 4

In [7]:

```
class AE4_DenoisingAutoencoder(nn.Module):
    def __init__(self):
        super(AE4_DenoisingAutoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784,256),
            nn.ReLU())
        self.decoder =nn.Sequential(
            nn.Linear(256,784),
            nn.ReLU())

    def forward(self,x):
        x=self.encoder(x)
        x=self.decoder(x)
        return x
```

In [8]:

```
def add_noise(img, noise_val):
    noise = torch.randn(img.size())*noise_val
    noisy_img = img + noise
    return noisy_img
```

In [18]:

```
#HYPERPARAMETERS
learning_rate=0.00008
epochs=10
batch_size= 300
```


In [10]:

```
model_4_a=AE4_DenoisingAutoencoder()  
criterion_4_a=nn.MSELoss()  
optimizer_4_a = torch.optim.Adam(model_4_a.parameters(),lr=learning_rate)  
  
training_loss_4_a=[]  
  
for epoch in range(epochs):  
    for images,labels in train_loader:  
        images=images.reshape(images.shape[0],-1)  
        noisy_images=add_noise(images,0.3)  
        outputs=model_4_a(noisy_images)  
        loss=criterion_4_a(outputs,images)  
        training_loss_4_a.append(loss.item())  
  
        optimizer_4_a.zero_grad()  
        loss.backward()  
        optimizer_4_a.step()  
  
    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed  
Epoch [ 2 / 10 ] : completed  
Epoch [ 3 / 10 ] : completed  
Epoch [ 4 / 10 ] : completed  
Epoch [ 5 / 10 ] : completed  
Epoch [ 6 / 10 ] : completed  
Epoch [ 7 / 10 ] : completed  
Epoch [ 8 / 10 ] : completed  
Epoch [ 9 / 10 ] : completed  
Epoch [ 10 / 10 ] : completed
```

In [12]:

```
model_4_b=AE4_DenoisingAutoencoder()  
criterion_4_b=nn.MSELoss()  
optimizer_4_b = torch.optim.Adam(model_4_b.parameters(),lr=learning_rate)  
  
training_loss_4_b=[]  
  
for epoch in range(epochs):  
    for images,labels in train_loader:  
        images=images.reshape(images.shape[0],-1)  
        noisy_images=add_noise(images,0.5)  
        outputs=model_4_b(noisy_images)  
        loss=criterion_4_b(outputs,images)  
        training_loss_4_b.append(loss.item())  
  
        optimizer_4_b.zero_grad()  
        loss.backward()  
        optimizer_4_b.step()  
  
    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed  
Epoch [ 2 / 10 ] : completed  
Epoch [ 3 / 10 ] : completed  
Epoch [ 4 / 10 ] : completed  
Epoch [ 5 / 10 ] : completed  
Epoch [ 6 / 10 ] : completed  
Epoch [ 7 / 10 ] : completed  
Epoch [ 8 / 10 ] : completed  
Epoch [ 9 / 10 ] : completed  
Epoch [ 10 / 10 ] : completed
```

In [13]:

```
model_4_c=AE4_DenoisingAutoencoder()  
criterion_4_c=nn.MSELoss()  
optimizer_4_c = torch.optim.Adam(model_4_c.parameters(),lr=learning_rate)  
  
training_loss_4_c=[]  
  
for epoch in range(epochs):  
    for images,labels in train_loader:  
        images=images.reshape(images.shape[0],-1)  
        noisy_images=add_noise(images,0.8)  
        outputs=model_4_c(noisy_images)  
        loss=criterion_4_c(outputs,images)  
        training_loss_4_c.append(loss.item())  
  
        optimizer_4_c.zero_grad()  
        loss.backward()  
        optimizer_4_c.step()  
  
    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed  
Epoch [ 2 / 10 ] : completed  
Epoch [ 3 / 10 ] : completed  
Epoch [ 4 / 10 ] : completed  
Epoch [ 5 / 10 ] : completed  
Epoch [ 6 / 10 ] : completed  
Epoch [ 7 / 10 ] : completed  
Epoch [ 8 / 10 ] : completed  
Epoch [ 9 / 10 ] : completed  
Epoch [ 10 / 10 ] : completed
```

In [14]:

```
model_4_d=AE4_DenoisingAutoencoder()
criterion_4_d=nn.MSELoss()
optimizer_4_d = torch.optim.Adam(model_4_d.parameters(),lr=learning_rate)

training_loss_4_d=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        images=images.reshape(images.shape[0],-1)
        noisy_images=add_noise(images,0.9)
        outputs=model_4_d(noisy_images)
        loss=criterion_4_d(outputs,images)
        training_loss_4_d.append(loss.item())

        optimizer_4_d.zero_grad()
        loss.backward()
        optimizer_4_d.step()


    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [15]:

```
plt.rcParams["figure.figsize"] = (15,6)

plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_a,label="noise factor
=0.3")
plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_b,label="noise factor
=0.5")
plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_c,label="noise factor
=0.8")
plt.plot(range(1,len(training_loss_4_a)+1),training_loss_4_d,label="noise factor
=0.9")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()
```



In [24]:

```
#INPUT WITH NOISE FACTOR=0.3,0.5,0.8,0.9 GIVEN TO VANILLA AE HIDLAYER=256
model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        # print(images.shape)
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.3)
        outputs_hid256_03,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.5)
        outputs_hid256_05,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.8)
        outputs_hid256_08,activations_hid256 = model_hid256(noisy_images.float())

model_hid256.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images=add_noise(images,0.9)
        outputs_hid256_09,activations_hid256 = model_hid256(noisy_images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)

ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
ax1.set_title('Original Image')
ax1.axis("off")

ax2.imshow(outputs_hid256_03[i].detach().numpy().reshape(28,28),cmap='gray')
ax2.set_title('AE_Reconst.Image_nf=0.3')
ax2.axis("off")

ax3.imshow(outputs_hid256_05[i].detach().numpy().reshape(28,28),cmap='gray')
ax3.set_title('AE_Reconst.Image_nf=0.5')
ax3.axis("off")

ax4.imshow(outputs_hid256_08[i].detach().numpy().reshape(28,28),cmap='gray')
ax4.set_title('AE_Reconst.Image_nf=0.8')
ax4.axis("off")

ax5.imshow(outputs_hid256_09[i].detach().numpy().reshape(28,28),cmap='gray')
ax5.set_title('AE_Reconst.Image_nf=0.9')
ax5.axis("off")

print("Reconstruction Error in VanillaAE with noise factor = 0.3 :",np.dot(((images[i].detach().numpy()/255.)-(outputs_hid256_03[i].detach().numpy()/255.)),((images[i].detach().numpy()/255.)-(outputs_hid256_03[i].detach().numpy()/255.)).T))
print("Reconstruction Error in VanillaAE with noise factor = 0.5 :",np.dot(((images[i].detach().numpy()/255.)-(outputs_hid256_05[i].detach().numpy()/255.)),((images[i].detach().numpy()/255.)-(outputs_hid256_05[i].detach().numpy()/255.)).T))
```

```
ages[i].detach().numpy()/255.)-(outputs_hid256_05[i].detach().numpy()/255.)).T))
print("Reconstruction Error in VanillaAE with noise factor = 0.8 :",np.dot(((ima
ges[i].detach().numpy()/255.)-(outputs_hid256_08[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_hid256_08[i].detach().numpy()/255.)).T))
print("Reconstruction Error in VanillaAE with noise factor = 0.9 :",np.dot(((ima
ges[i].detach().numpy()/255.)-(outputs_hid256_09[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_hid256_09[i].detach().numpy()/255.)).T))
```

Reconstruction Error in VanillaAE with noise factor = 0.3 : 12.29054
5043085082
Reconstruction Error in VanillaAE with noise factor = 0.5 : 12.28854
7516879056
Reconstruction Error in VanillaAE with noise factor = 0.8 : 12.31252
712926534
Reconstruction Error in VanillaAE with noise factor = 0.9 : 12.32003
4453259053



In [16]:

```

model_4_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.3)
        outputs_4_a = model_4_a(noisy_images.float())

model_4_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.5)
        outputs_4_b = model_4_b(noisy_images.float())

model_4_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.8)
        outputs_4_c = model_4_c(noisy_images.float())

model_4_d.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.9)
        outputs_4_d = model_4_d(noisy_images.float())

plt.rcParams["figure.figsize"] = (15,6)
i=5
if i==5:
    fig, (ax1, ax2, ax3, ax4, ax5) = plt.subplots(1,5)

    ax1.imshow(images[i].detach().numpy().reshape(28,28), cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_4_a[i].detach().numpy().reshape(28,28), cmap='gray')
    ax2.set_title('Reconstruction: nf=0.3')
    ax2.axis("off")

    ax3.imshow(outputs_4_b[i].detach().numpy().reshape(28,28), cmap='gray')
    ax3.set_title('Reconstruction: nf=0.5')
    ax3.axis("off")

    ax4.imshow(outputs_4_c[i].detach().numpy().reshape(28,28), cmap='gray')
    ax4.set_title('Reconstruction: nf=0.8')
    ax4.axis("off")

    ax5.imshow(outputs_4_d[i].detach().numpy().reshape(28,28), cmap='gray')
    ax5.set_title('Reconstruction: nf=0.9')
    ax5.axis("off")

    plt.show()

    print("Reconstruction Error in DenoisingAE with noise factor = 0.3 :", np.dot
    (((images[i].detach().numpy()/255.)-(outputs_4_a[i].detach().numpy()/255.)), ((im
    ages[i].detach().numpy()/255.)-(outputs_4_a[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.5 :", np.dot


```



```

(((images[i].detach().numpy()/255.)-(outputs_4_b[i].detach().numpy()/255.)),((images[i].detach().numpy()/255.)-(outputs_4_b[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.8 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_c[i].detach().numpy()/255.)),((images[i].detach().numpy()/255.)-(outputs_4_c[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.9 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_d[i].detach().numpy()/255.)),((images[i].detach().numpy()/255.)-(outputs_4_d[i].detach().numpy()/255.)).T))
    print("_____")

```



```

Reconstruction Error in DenoisingAE with noise factor = 0.3 : 11.504
991739148307
Reconstruction Error in DenoisingAE with noise factor = 0.5 : 14.426
701488962669
Reconstruction Error in DenoisingAE with noise factor = 0.8 : 19.260
798583259362
Reconstruction Error in DenoisingAE with noise factor = 0.9 : 21.399
583871015913

```

In [25]:

```

#DENOISING ENCODER TRAINED ON NOISE LEVEL=0.3 AND TESTED RECONSTRUCTION FOR NOIS
E LEVEL 0.3,0.5,0.8,0.9
model_4_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images = images.reshape(10,28*28)
        noisy_images = add_noise(images,0.3)
        outputs_4_a = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.5)
        outputs_4_b = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.8)
        outputs_4_c = model_4_a(noisy_images.float())
        noisy_images = add_noise(images,0.9)
        outputs_4_d = model_4_a(noisy_images.float())
i=5
if i==5:
    fig, (ax1, ax2,ax3,ax4,ax5) = plt.subplots(1,5)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_4_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Reconstruction: nf=0.3')
    ax2.axis("off")

    ax3.imshow(outputs_4_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Reconstruction: nf=0.5')
    ax3.axis("off")


    ax4.imshow(outputs_4_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Reconstruction: nf=0.8')
    ax4.axis("off")

    ax5.imshow(outputs_4_d[i].detach().numpy().reshape(28,28),cmap='gray')
    ax5.set_title('Reconstruction: nf=0.9')
    ax5.axis("off")

    plt.show()

    print("Reconstruction Error in DenoisingAE with noise factor = 0.3 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_a[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_4_a[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.5 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_b[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_4_b[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.8 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_c[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_4_c[i].detach().numpy()/255.)).T))
    print("Reconstruction Error in DenoisingAE with noise factor = 0.9 :",np.dot
(((images[i].detach().numpy()/255.)-(outputs_4_d[i].detach().numpy()/255.)),((im
ages[i].detach().numpy()/255.)-(outputs_4_d[i].detach().numpy()/255.)).T))
    print("_____")

```



Reconstruction Error in DenoisingAE with noise factor = 0.3 : 11.504407994823417

Reconstruction Error in DenoisingAE with noise factor = 0.5 : 11.499836309490743

Reconstruction Error in DenoisingAE with noise factor = 0.8 : 11.495695934586088

Reconstruction Error in DenoisingAE with noise factor = 0.9 : 11.50995439903982

In [23]:

```
##VISUALISING ENCODER WEIGHTS
```

```
plt.rcParams["figure.figsize"] = (15,6)
```

```
fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)
```

```
ax1.imshow(model_hid256.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
```

```
ax1.set_title('VaniallaAE')
```

```
ax1.axis("off")
```

```
ax2.imshow(model_4_a.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
```

```
ax2.set_title('DenoisingAE nf=0.3')
```

```
ax2.axis("off")
```

```
ax3.imshow(model_4_b.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
```

```
ax3.set_title('DenoisingAE nf=0.5')
```


```
ax3.axis("off")
```

```
ax4.imshow(model_4_c.encoder[0].weight.detach().numpy()[0].reshape(28,28),cmap='gray')
```

```
ax4.set_title('DenoisingAE nf=0.8')
```

```
ax4.axis("off")
```

```
plt.show()
```



Observations:

- Vanilla autoencoder when given corrupted input has reconstruction error more than Denoising Autoencoder trained on noise level 0.3.
- Though not very clearly visible but encoder weights of denoising AE have some penstroke detector type visualisation.

Question 5

In [26]:

```

class AE5_ConvAE_with_unpooling(nn.Module): #define unpooling outside the decoder and separately in forward nn.Sequential just takes one input

    def __init__(self): #class constructor
        super(AE5_ConvAE_with_unpooling,self).__init__() #calls the parent constructor

        #initializing the encoder module
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        ) # 28x28x1 to 14x14x8
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        ) #14x14x8 to 7x7x16
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        ) #7x7x16 to 3x3x16

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(nn.Identity()) #7x7x16 to 7x7x16
        self.decoder_conv2 = nn.Sequential(
            nn.Conv2d(16,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU()
        ) #14x14x16 to 14x14x8
        self.decoder_conv3 = nn.Sequential(
            nn.Conv2d(8,1, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU()
        ) #28x28x8 to 28x28x1

        #defining the unpooling operation
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x): #defines the forward pass and also the structure of the network thus helping backprop

        encoded_input,indices1 = self.encoder_conv1(x.float()) # 28x28x1 to 14x14x8
        encoded_input,indices2 = self.encoder_conv2(encoded_input) #14x14x8 to 7x7x16
        encoded_input,indices3 = self.encoder_conv3(encoded_input) #7x7x16 to 3x3x16

        reconstructed_input = self.unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7, 7])) #3x3x16 to 7x7x16
        reconstructed_input = self.decoder_conv1(reconstructed_input) #7x7x16 to 7x7x16
        reconstructed_input = self.unpool(reconstructed_input,indices2) #7x7x16 to 14x14x16
        reconstructed_input = self.decoder_conv2(reconstructed_input) #14x14x16 to 14x14x8
        reconstructed_input = self.unpool(reconstructed_input,indices1) #14x14x8 to 28x28x8

```

```

        reconstructed_input      = self.decoder_conv3(reconstructed_input)#28x28x
6 to 28x28x1

    return reconstructed_input,encoded_input

```

In [27]:

```

model_5_a = AE5_ConvAE_with_unpooling()
criterion_5_a = nn.MSELoss()
optimizer_5_a = torch.optim.Adam(model_5_a.parameters(),lr=0.001)

training_loss_5_a=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs,_=model_5_a(images)
        loss=criterion_5_a(outputs,images)
        training_loss_5_a.append(loss.item())

        optimizer_5_a.zero_grad()
        loss.backward()
        optimizer_5_a.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")

```

```

Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed

```

In [28]:

```

class AE5_ConvAE_with_deconv(nn.Module):
    def __init__(self):
        super(AE5_ConvAE_with_deconv,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )

        #decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 2),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )

    def forward(self,x):

        encoded_input  = self.encoder_conv1(x.float())
        encoded_input  = self.encoder_conv2(encoded_input)
        encoded_input  = self.encoder_conv3(encoded_input)

        reconstructed_input = self.decoder_conv1(encoded_input)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
        reconstructed_input = self.decoder_conv3(reconstructed_input)

        return reconstructed_input,encoded_input

```

In [29]:

```
model_5_b = AE5_ConvAE_with_deconv()
criterion_5_b = nn.MSELoss()
optimizer_5_b = torch.optim.Adam(model_5_b.parameters(),lr=0.001)

training_loss_5_b=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs,_=model_5_b(images)
        loss=criterion_5_b(outputs,images)
        training_loss_5_b.append(loss.item())

        optimizer_5_b.zero_grad()
        loss.backward()
        optimizer_5_b.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```


In [30]:

```

class AE5_ConvAE_with_deconv_unpool(nn.Module):
    def __init__(self):
        super(AE5_ConvAE_with_deconv_unpool,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )

        #unpooling
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x): #defines the forward pass and also the structure of the
network thus helping backprop

        encoded_input,indices1 = self.encoder_conv1(x.float())
        encoded_input,indices2 = self.encoder_conv2(encoded_input)
        encoded_input,indices3 = self.encoder_conv3(encoded_input)

        reconstructed_input = self.unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7, 7]))
        reconstructed_input = self.decoder_conv1(reconstructed_input)
        reconstructed_input = self.unpool(reconstructed_input,indices2)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
        reconstructed_input = self.unpool(reconstructed_input,indices1)
        reconstructed_input = self.decoder_conv3(reconstructed_input)

        return reconstructed_input,encoded_input

```

In [31]:

```
model_5_c = AE5_ConvAE_with_deconv_unpool()
criterion_5_c = nn.MSELoss()
optimizer_5_c = torch.optim.Adam(model_5_c.parameters(),lr=0.001)

training_loss_5_c=[]

for epoch in range(epochs):
    for images,labels in train_loader:
        outputs,_=model_5_c(images)
        loss=criterion_5_c(outputs,images)
        training_loss_5_c.append(loss.item())


        optimizer_5_c.zero_grad()
        loss.backward()
        optimizer_5_c.step()

    print("Epoch", "[", epoch+1, "/", epochs, "]", ": completed")
```

```
Epoch [ 1 / 10 ] : completed
Epoch [ 2 / 10 ] : completed
Epoch [ 3 / 10 ] : completed
Epoch [ 4 / 10 ] : completed
Epoch [ 5 / 10 ] : completed
Epoch [ 6 / 10 ] : completed
Epoch [ 7 / 10 ] : completed
Epoch [ 8 / 10 ] : completed
Epoch [ 9 / 10 ] : completed
Epoch [ 10 / 10 ] : completed
```

In [32]:

```
plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_a,label="ConvAE Unpooling")
plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_b,label="ConvAE Deconv")
plt.plot(range(1,len(training_loss_5_a)+1),training_loss_5_c,label="ConvAE Deconv+Unpooling")
plt.legend()
plt.grid()
plt.title("Training Loss Vs Iterations")
plt.xlabel("Number of Iterations")
plt.ylabel("Training Loss")
plt.show()
```



In [33]:

```

model_5_a.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images=images.reshape(10,1,28,28)
        outputs_5_a,_ = model_5_a(images.float())

model_5_b.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images=images.reshape(10,1,28,28)
        outputs_5_b,_ = model_5_b(images.float())

activation_5_c=[]
model_5_c.eval()
with torch.no_grad():
    for images in test_sample_loader:
        images=images.reshape(10,1,28,28)
        outputs_5_c,_ = model_5_c(images.float())

plt.rcParams["figure.figsize"] = (15,6)

for i in range(10):
    fig, (ax1, ax2,ax3,ax4) = plt.subplots(1,4)

    ax1.imshow(images[i].detach().numpy().reshape(28,28),cmap='gray')
    ax1.set_title('Original Image')
    ax1.axis("off")

    ax2.imshow(outputs_5_a[i].detach().numpy().reshape(28,28),cmap='gray')
    ax2.set_title('Reconstruction: CAE_Unpooling')
    ax2.axis("off")


    ax3.imshow(outputs_5_b[i].detach().numpy().reshape(28,28),cmap='gray')
    ax3.set_title('Reconstruction: CAE_Deconv')
    ax3.axis("off")

    ax4.imshow(outputs_5_c[i].detach().numpy().reshape(28,28),cmap='gray')
    ax4.set_title('Reconstruction: CAE_Deconv+Unpooling')
    ax4.axis("off")

    plt.show()

    print("Reconstruction Error in ConvAE Unpooling:",np.sum(np.dot(((images[i].de
tach().numpy()/255)-(outputs_5_a[i].detach().numpy()/255)),((images[i].detach().
numpy()/255)-(outputs_5_a[i].detach().numpy()/255)).T)))
    print("Reconstruction Error in ConvAE Deconv:",np.sum(np.dot(((images[i].detac
h().numpy()/255)-(outputs_5_b[i].detach().numpy()/255)),((images[i].detach().num
py()/255)-(outputs_5_b[i].detach().numpy()/255)).T)))
    print("Reconstruction Error in ConvAE Deconv+Unpooling:",np.sum(np.dot(((image
s[i].detach().numpy()/255)-(outputs_5_c[i].detach().numpy()/255)),((images[i].de
tach().numpy()/255)-(outputs_5_c[i].detach().numpy()/255)).T)))
    print("_____")


```



Reconstruction Error in ConvAE Unpooling: 1.7310008604428355

Reconstruction Error in ConvAE Deconv: 26.001679380843555


Reconstruction Error in ConvAE Deconv+Unpooling: 11.89297063287778



Reconstruction Error in ConvAE Unpooling: 36.132811210593644

Reconstruction Error in ConvAE Deconv: 118.04138416590519


Reconstruction Error in ConvAE Deconv+Unpooling: 93.50976356860426




Reconstruction Error in ConvAE Unpooling: 0.2794066287723087

Reconstruction Error in ConvAE Deconv: 130.92309175297225


Reconstruction Error in ConvAE Deconv+Unpooling: 56.26377338926373




Reconstruction Error in ConvAE Unpooling: 21.321372713953124
 Reconstruction Error in ConvAE Deconv: 95.42958725539685
 Reconstruction Error in ConvAE Deconv+Unpooling: 38.815502513706875




Reconstruction Error in ConvAE Unpooling: 6.22126437500681
 Reconstruction Error in ConvAE Deconv: 61.48894311018323
 Reconstruction Error in ConvAE Deconv+Unpooling: 9.12990892796395




Reconstruction Error in ConvAE Unpooling: 71.43498495324837
 Reconstruction Error in ConvAE Deconv: 212.74946483973974
 Reconstruction Error in ConvAE Deconv+Unpooling: 168.52247765694244




Reconstruction Error in ConvAE Unpooling: 25.535638771069415
 Reconstruction Error in ConvAE Deconv: 95.00550731184873
 Reconstruction Error in ConvAE Deconv+Unpooling: 35.801245134441594



Reconstruction Error in ConvAE Unpooling: 29.760346086314858
 Reconstruction Error in ConvAE Deconv: 101.14872994906563
 Reconstruction Error in ConvAE Deconv+Unpooling: 53.518502795113626



Reconstruction Error in ConvAE Unpooling: 24.24100375371164
 Reconstruction Error in ConvAE Deconv: 86.583618580179
 Reconstruction Error in ConvAE Deconv+Unpooling: 41.803233553557405



Reconstruction Error in ConvAE Unpooling: 41.52189672531831
 Reconstruction Error in ConvAE Deconv: 205.93324026718437
 Reconstruction Error in ConvAE Deconv+Unpooling: 80.49537718296492

In [34]:

```
model_5_a.encoder_conv1[0].weight.detach().numpy().squeeze().shape
```

Out[34]:

(8, 3, 3)

In [35]:

```
#Function for visualisation of weights
from torchvision import utils

def visTensor(tensor, ch=0, allkernels=False, nrow=8, padding=1):
    n,c,w,h = tensor.shape

    if allkernels: tensor = tensor.view(n*c, -1, w, h)
    elif c != 3: tensor = tensor[:,ch,:].unsqueeze(dim=1)

    rows = np.min((tensor.shape[0] // nrow + 1, 64))
    grid = utils.make_grid(tensor, nrow=nrow, normalize=True, padding=padding)
    plt.figure( figsize=(nrow,rows) )
    plt.imshow(grid.numpy().transpose((1, 2, 0)))
```

In [36]:

```
#VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH UNPOOLING
```

```
filter = model_5_a.decoder_conv2[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv2 Weights')  
plt.show()
```

```
filter = model_5_a.decoder_conv3[0].weight.data.clone()  
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')  
plt.ioff()  
plt.title('decoder_conv3 Weights')  
plt.show()
```

decoder_conv2 Weights



decoder_conv3 Weights



In [37]:

#VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH DECONVOLUTION

```
filter = model_5_b.decoder_conv1[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv1 Weights')
plt.show()
```


```
filter = model_5_b.decoder_conv2[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv2 Weights')
plt.show()
```


```
filter = model_5_b.decoder_conv3[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)
```

```
plt.axis('off')
plt.ioff()
plt.title('decoder_conv3 Weights')
plt.show()
```


decoder_conv1 Weights



decoder_conv2 Weights



decoder_conv3 Weights



In [38]:

```
#VISUALISING DECODER WEIGHTS FOR CONVOLUTION AUTOENCODER WITH DECONVOLUTION+UNPOOLING

filter = model_5_c.decoder_conv1[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv1 Weights')
plt.show()


filter = model_5_c.decoder_conv2[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv2 Weights')
plt.show()


filter = model_5_c.decoder_conv3[0].weight.data.clone()
visTensor(filter, ch=0, allkernels=False)

plt.axis('off')
plt.ioff()
plt.title('decoder_conv3 Weights')
plt.show()
```

decoder_conv1 Weights



decoder_conv2 Weights



decoder_conv3 Weights



Observations:

- Decoder weights of Deconvolution with Unpooling and Unpooling are smaller than the only Deconvolution one.
- By looking at reconstruction error, it looks like unpooling and deconv+ unpooling does better on reconstruction of images.
- Visually, reconstructed images using Deconvolution are appealing.