

Importing Required Libraries

In [151]:

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
from torch.utils.data.dataset import TensorDataset
from torch.utils.data import DataLoader
```

DATASET LOADING AND PREPARATION

In [152]:

```
train_dataset = torchvision.datasets.MNIST(root='./data', train=True, transform=
transforms.ToTensor(), download=True)
test_dataset = torchvision.datasets.MNIST(root='./data', train=False, transform=
transforms.ToTensor())
```

In [153]:

```
loader_train = torch.utils.data.DataLoader(dataset=train_dataset, batch_size = le
n(train_dataset) ,shuffle=True)
loader_test = torch.utils.data.DataLoader(dataset=test_dataset, batch_size = len(
test_dataset) ,shuffle=True)
```

In [154]:

```
X_train,y_train=next(iter(loader_train))
X_test,y_test=next(iter(loader_test))
```

In [155]:

```
X_train=X_train.numpy()
y_train=y_train.numpy()
X_test=X_test.numpy()
y_test=y_test.numpy()
```

In [156]:

```
X_train_flattened=X_train.reshape(X_train.shape[0],X_train.shape[1]*X_train.shap
e[2]*X_train.shape[3])
X_test_flattened=X_test.reshape(X_test.shape[0],X_test.shape[1]*X_test.shape[2]*
X_test.shape[3])
```

One-Hot Encoding

In [157]:

```
train_labels_encoded = []
for i in y_train:
    A=np.array([0]*10)
    A[i]=1
    train_labels_encoded.append(A)
y_train_encoded=np.array(train_labels_encoded)
```

In [158]:

```
test_labels_encoded = []
for i in y_test:
    A=np.array([0]*10)
    A[i]=1
    test_labels_encoded.append(A)
y_test_encoded=np.array(test_labels_encoded)
```

In [159]:

```
X_train_flattened_torch=torch.from_numpy(X_train_flattened)
y_train_encoded_torch=torch.from_numpy(y_train_encoded)
```

In [160]:

```
X_test_flattened_torch=torch.from_numpy(X_test_flattened)
y_test_encoded_torch=torch.from_numpy(y_test_encoded)
```

Here we have prepared test and train dataset with flattened images and one hot encoded labels

In [161]:

```
Train_Dataset=TensorDataset(X_train_flattened_torch,y_train_encoded_torch)
```

In [162]:

```
Test_Dataset=TensorDataset(X_test_flattened_torch,y_test_encoded_torch)
```

HELPER FUNCTIONS

In [163]:

```
def relu(z):
    return np.maximum(0,z)

def softmax(z):
    return np.exp(z)/sum(np.exp(z))
```

In [164]:

```
def dif_relu(z):
    return z>0
```

In [165]:

```
def glorot_initialisation(output_n,input_n):
    M=np.sqrt(6/(input_n+output_n))
    W=np.random.uniform(low=-M, high=M, size=(output_n,input_n))
    b=np.random.uniform(low=-M, high=M, size=(output_n,1))
    return W,b
```

In [166]:

```
def initialize_parameters(layer_dims):
    """
    Arguments:
        layer_dims -- python array (list) containing the dimensions of each layer in
our network

    Returns:
        parameters -- python dictionary containing your parameters "W1", "b1", ...,
"WL", "bL":
            W1 -- weight matrix of shape (layer_dims[1], layer_dims[1-
1])
            b1 -- bias vector of shape (layer_dims[1], 1)
            Wl -- weight matrix of shape (layer_dims[l-1], layer_dims
[l])
            bl -- bias vector of shape (1, layer_dims[l])

    Tips:
        - For example: the layer_dims for the "Planar Data classification model" wou
ld have been [2,2,1].
        This means W1's shape was (2,2), b1 was (1,2), W2 was (2,1) and b2 was (1,
1). Now you have to generalize it!
        - In the for loop, use parameters['W' + str(l)] to access Wl, where l is the
iterative integer.
    """

    np.random.seed(1390)
    parameters = {}
    L = len(layer_dims) # number of layers in the network

    for l in range(1, L):

        # M=np.sqrt(6/(self.input_n+self.output_n))
        parameters['W' + str(l)],parameters['b' + str(l)] = glorot_initialisati
on(layer_dims[l], layer_dims[l-1])

        assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l-1
]))
        assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

    return parameters
```

In [167]:

```

def forward_propagation(X, parameters):
    """
    SHAPE OF X = 784,samples(i.e. 64 for a batch)
    Implements the forward propagation (and computes the loss) presented in Figure 2.

    Arguments:
    X -- input dataset, of shape (input size, number of examples)
    parameters -- python dictionary containing your parameters "W1", "b1", "W2", "b2", "W3", "b3", "W4", "b4":
        W1 -- weight matrix of shape (500,784)
        b1 -- bias vector of shape (500,1)
        W2 -- weight matrix of shape (250,500)
        b2 -- bias vector of shape (250,1)
        W3 -- weight matrix of shape (100,250)
        b3 -- bias vector of shape (100,1)
        W4 -- weight matrix of shape (10,100)
        b4 -- bias vector of shape (10,1)

    Returns:
    loss -- the loss function (vanilla logistic loss)
    """

    cache={}

    cache['Z2']=np.dot(parameters['W1'],X)+parameters['b1']
    cache['A2']=relu(cache['Z2'])

    cache['Z3']=np.dot(parameters['W2'],cache['A2'])+parameters['b2']
    cache['A3']=relu(cache['Z3'])

    cache['Z4']=np.dot(parameters['W3'],cache['A3'])+parameters['b3']
    cache['A4']=relu(cache['Z4'])

    cache['Z5']=np.dot(parameters['W4'],cache['A4'])+parameters['b4']
    cache['A5']=softmax(cache['Z5'])

    A_last=cache['A5']

    cache['W1']=parameters['W1']
    cache['b1']=parameters['b1']

    cache['W2']=parameters['W2']
    cache['b2']=parameters['b2']

    cache['W3']=parameters['W3']
    cache['b3']=parameters['b3']

    cache['W4']=parameters['W4']
    cache['b4']=parameters['b4']

    # cache = {"W1": W1, "b1": b1, "Z2": Z2, "A2": A2,
    #          "W2": W2, "b2": b2, "Z3": Z3, "A3": A3,
    #          "W3": W3, "b3": b3, "Z4": Z4, "A4": A4,
    #          "W4": W4, "b4": b4, "Z5": Z5, "A5": A5}

```

```
return A_last, cache
```

In [168]:

```
# Calculating the loss function using the cross entropy
"""Arguments:
    A -- post-activation, output of forward propagation
    Y -- "true" labels vector, same shape as A

Returns:
    cost - value of the cost function
"""
def compute_cost(A, Y):
    #A is predicted
    #Y is actual
    m = Y.shape[1]
    logprobs = np.multiply(-np.log(A), Y) + np.multiply(-np.log(1 - A), 1 - Y)
    cost = 1./m * np.nansum(logprobs)
    return cost
```

In [169]:

```
def backward_propagation(X, Y, cache):

    m=X.shape[1]

    grads={}

    grads['dz5']=cache['A5']-Y
    grads['dw4']= 1./m * np.dot(grads['dz5'],cache['A4'].T)
    grads['db4']= 1./m * np.sum(grads['dz5'],axis=1,keepdims=True)

    grads['dA4']=np.dot(cache['W4'].T,grads['dz5'])
    grads['dz4']=np.multiply(grads['dA4'],dif_relu(cache['Z4']))
    grads['dw3']=1./m * np.dot(grads['dz4'],cache['A3'].T)
    grads['db3']=1./m * np.sum(grads['dz4'],axis=1,keepdims=True)

    grads['dA3']=np.dot(cache['W3'].T,grads['dz4'])
    grads['dz3']=np.multiply(grads['dA3'],dif_relu(cache['Z3']))
    grads['dw2']=1./m * np.dot(grads['dz3'],cache['A2'].T)
    grads['db2']=1./m * np.sum(grads['dz3'],axis=1,keepdims=True)

    grads['dA2']=np.dot(cache['W2'].T,grads['dz3'])
    grads['dz2']=np.multiply(grads['dA2'],dif_relu(cache['Z2']))
    grads['dw1']=1./m * np.dot(grads['dz2'],X.T)
    grads['db1']=1./m * np.sum(grads['dz2'],axis=1,keepdims=True)

    return grads
```

In [170]:

```
def update_parameters(parameters, grads, learning_rate):

    updated_parameters={}

    updated_parameters['W1']=parameters['W1']-learning_rate*grads['dW1']
    updated_parameters['b1']=parameters['b1']-learning_rate*grads['db1']

    updated_parameters['W2']=parameters['W2']-learning_rate*grads['dW2']
    updated_parameters['b2']=parameters['b2']-learning_rate*grads['db2']

    updated_parameters['W3']=parameters['W3']-learning_rate*grads['dW3']
    updated_parameters['b3']=parameters['b3']-learning_rate*grads['db3']

    updated_parameters['W4']=parameters['W4']-learning_rate*grads['dW4']
    updated_parameters['b4']=parameters['b4']-learning_rate*grads['db4']

    return updated_parameters
```

In [171]:

```
#Calculating accuracy of the parameter at the output
"""
    Arguments:
    y_actual - given in the dataset / also called as the ground truth
    y_pred - generated from the neural network , after a series of forward and b
ackprop

    Returns:
    accuracy = finding the matches of the predicted vs the actual
"""
def calculate_accuracy(y_actual,y_pred):
    accuracy = np.count_nonzero(np.argmax(y_pred,axis=0)==np.argmax(y_actual,axis=0))/y_actual.shape[1]
    return accuracy
```

In [172]:

```
def predict(X,Y,parameters):

    """
    This function is used to predict the results of a n-layer neural network.

    Arguments:
    X -- data set of examples you would like to label
    Y -- data set of examples
    parameters -- parameters of the trained model

    Returns:
    ypred -- predictions for the given dataset X
    """

    y_pred,cache=forward_propagation(X,parameters)
    return y_pred
```

In [173]:

```
def model(Train_Dataset,layer_dimensions,total_epochs=15,Batch_Size=64,learning_
rate=0.01):

    costs=[]
    accuracy=[]

    parameters=initialize_parameters(layer_dimensions)
    num_iterations=len(Train_Dataset)//Batch_Size

    #Train_Dataset=TensorDataset(X_training,Y_training)
    for epoch in range(total_epochs):
        for iteration in range(num_iterations):
            Data_Loader=torch.utils.data.DataLoader(dataset=Train_Dataset,batch_size=6
4, shuffle=True)

            data_iter=iter(Data_Loader)
            Data=next(data_iter)
            X,y=Data #X.shape=(batch_size,784) y.shape=(batch_size,10)
            X=X.numpy()
            y=y.numpy()
            a5,cache=forward_propagation(X.T,parameters)
            cost=compute_cost(a5,y.T)
            gradients=backward_propagation(X.T,y.T,cache)
            parameters=update_parameters(parameters,gradients,learning_rate)

            if iteration%200==0:
                print("epoch: ",epoch+1,"/",total_epochs, "  iteration= ",iteration+1,
"/",num_iterations, "  Loss: ",cost)
                accuracy.append(calculate_accuracy(y.T,a5))
                costs.append(cost)

    return accuracy, costs, parameters
```

In [174]:

```
layer_dimensions=[784,500,250,100,10]  
Train_accuracy,Train_costs,Trained_parameters=model(Train_Dataset,layer_dimensions,15,64,0.01)
```



```
epoch: 1 / 15 iteration= 1 / 937 Loss: 3.29407977813217
epoch: 1 / 15 iteration= 201 / 937 Loss: 1.6627210071024752
epoch: 1 / 15 iteration= 401 / 937 Loss: 0.9182173362911166
epoch: 1 / 15 iteration= 601 / 937 Loss: 0.8881546656754662
epoch: 1 / 15 iteration= 801 / 937 Loss: 0.6482982901984773
epoch: 2 / 15 iteration= 1 / 937 Loss: 0.7209205518947087
epoch: 2 / 15 iteration= 201 / 937 Loss: 0.724783929373435
epoch: 2 / 15 iteration= 401 / 937 Loss: 0.6026104563690713
epoch: 2 / 15 iteration= 601 / 937 Loss: 0.564973011083954
epoch: 2 / 15 iteration= 801 / 937 Loss: 0.5325791780407113
epoch: 3 / 15 iteration= 1 / 937 Loss: 0.3174605385417676
epoch: 3 / 15 iteration= 201 / 937 Loss: 0.4000062064141455
epoch: 3 / 15 iteration= 401 / 937 Loss: 0.36046783529714765
epoch: 3 / 15 iteration= 601 / 937 Loss: 0.5979444373404367
epoch: 3 / 15 iteration= 801 / 937 Loss: 0.4656260592292043
epoch: 4 / 15 iteration= 1 / 937 Loss: 0.29786494420348464
epoch: 4 / 15 iteration= 201 / 937 Loss: 0.1809752282333488
epoch: 4 / 15 iteration= 401 / 937 Loss: 0.5251086449164396
epoch: 4 / 15 iteration= 601 / 937 Loss: 0.2872652552504207
epoch: 4 / 15 iteration= 801 / 937 Loss: 0.3186067204964655
epoch: 5 / 15 iteration= 1 / 937 Loss: 0.18791682238315255
epoch: 5 / 15 iteration= 201 / 937 Loss: 0.42774623406279555
epoch: 5 / 15 iteration= 401 / 937 Loss: 0.30644336369836855
epoch: 5 / 15 iteration= 601 / 937 Loss: 0.2800557165951709
epoch: 5 / 15 iteration= 801 / 937 Loss: 0.3624913244213082
epoch: 6 / 15 iteration= 1 / 937 Loss: 0.23434464721735915
epoch: 6 / 15 iteration= 201 / 937 Loss: 0.23342689514633796
epoch: 6 / 15 iteration= 401 / 937 Loss: 0.2057274802447116
epoch: 6 / 15 iteration= 601 / 937 Loss: 0.1385572115149952
epoch: 6 / 15 iteration= 801 / 937 Loss: 0.19863783561988807
epoch: 7 / 15 iteration= 1 / 937 Loss: 0.26699711227846823
epoch: 7 / 15 iteration= 201 / 937 Loss: 0.18235101307197593
epoch: 7 / 15 iteration= 401 / 937 Loss: 0.2152246770143896
epoch: 7 / 15 iteration= 601 / 937 Loss: 0.29614714421607763
epoch: 7 / 15 iteration= 801 / 937 Loss: 0.19705903820443493
epoch: 8 / 15 iteration= 1 / 937 Loss: 0.27326279637596906
epoch: 8 / 15 iteration= 201 / 937 Loss: 0.4599160344284716
epoch: 8 / 15 iteration= 401 / 937 Loss: 0.32950710960831475
epoch: 8 / 15 iteration= 601 / 937 Loss: 0.1638453647801498
epoch: 8 / 15 iteration= 801 / 937 Loss: 0.18328637774679601
epoch: 9 / 15 iteration= 1 / 937 Loss: 0.2585832367096489
epoch: 9 / 15 iteration= 201 / 937 Loss: 0.21172953308772044
epoch: 9 / 15 iteration= 401 / 937 Loss: 0.1611034559981763
epoch: 9 / 15 iteration= 601 / 937 Loss: 0.2459757621937695
epoch: 9 / 15 iteration= 801 / 937 Loss: 0.20608936146058907
epoch: 10 / 15 iteration= 1 / 937 Loss: 0.2598337096238988
epoch: 10 / 15 iteration= 201 / 937 Loss: 0.1687004532876285
epoch: 10 / 15 iteration= 401 / 937 Loss: 0.18983107423201026
epoch: 10 / 15 iteration= 601 / 937 Loss: 0.1506146043655602
epoch: 10 / 15 iteration= 801 / 937 Loss: 0.1610571453690105
epoch: 11 / 15 iteration= 1 / 937 Loss: 0.09806433829943666
epoch: 11 / 15 iteration= 201 / 937 Loss: 0.2171259401366003
epoch: 11 / 15 iteration= 401 / 937 Loss: 0.09400254651080858
epoch: 11 / 15 iteration= 601 / 937 Loss: 0.2631639165233999
epoch: 11 / 15 iteration= 801 / 937 Loss: 0.26664382135037296
epoch: 12 / 15 iteration= 1 / 937 Loss: 0.21250345051747183
epoch: 12 / 15 iteration= 201 / 937 Loss: 0.19633486965936475
epoch: 12 / 15 iteration= 401 / 937 Loss: 0.31763841942343374
epoch: 12 / 15 iteration= 601 / 937 Loss: 0.23522792415028237
epoch: 12 / 15 iteration= 801 / 937 Loss: 0.05559882672794416
epoch: 13 / 15 iteration= 1 / 937 Loss: 0.11780774461068838
```


```
epoch: 13 / 15   iteration= 201 / 937   Loss:  0.20408612938846155
epoch: 13 / 15   iteration= 401 / 937   Loss:  0.17027088492964992
epoch: 13 / 15   iteration= 601 / 937   Loss:  0.18728915671640933
epoch: 13 / 15   iteration= 801 / 937   Loss:  0.04128862424173241
epoch: 14 / 15   iteration= 1 / 937     Loss:  0.07128548750104899
epoch: 14 / 15   iteration= 201 / 937   Loss:  0.09600554456200874
epoch: 14 / 15   iteration= 401 / 937   Loss:  0.1138127788763494
epoch: 14 / 15   iteration= 601 / 937   Loss:  0.3104434510976468
epoch: 14 / 15   iteration= 801 / 937   Loss:  0.09724889816207247
epoch: 15 / 15   iteration= 1 / 937     Loss:  0.07995103685920389
epoch: 15 / 15   iteration= 201 / 937   Loss:  0.18837370122730704
epoch: 15 / 15   iteration= 401 / 937   Loss:  0.08159515253388488
epoch: 15 / 15   iteration= 601 / 937   Loss:  0.15917562031532834
epoch: 15 / 15   iteration= 801 / 937   Loss:  0.1460621810813611
```

In [175]:

```
plt.plot(Train_costs)
```

Out[175]:

[<matplotlib.lines.Line2D at 0x7f71cb356fd0>]




In [176]:

```
plt.plot(Train_accuracy)
```

Out[176]:

[<matplotlib.lines.Line2D at 0x7f71e1e96510>]



In [177]:

```
Train_accuracy[-1]
```

Out[177]:

0.984375

In [178]:

```
Test_Data_Loader=torch.utils.data.DataLoader(dataset=Test_Dataset,batch_size=len
(Test_Dataset),shuffle=True)
data_iter=iter(Test_Data_Loader)
Test_Data=next(data_iter)
X,y=Test_Data
X=X.numpy()
y=y.numpy()
y_predicted,cache_out=forward_propagation(X.T,Trained_parameters)
cost=compute_cost(y_predicted,y.T)
Test_accuracy=calculate_accuracy(y.T,y_predicted)
```

In [179]:

```
Test_accuracy
```

Out[179]:

0.9712

In [180]:

```
Y_Predicted=np.array(np.argmax(y_predicted,axis=0))
```

In [181]:

```
Y_Actual=np.array(np.argmax(y.T,axis=0) )
```

CONFUSION MATRIX AND CLASSIFICATION REPORT

In [182]:


```
from sklearn.metrics import confusion_matrix
import seaborn as sns
plt.figure(figsize=(10,10))
conf_matrix = (confusion_matrix(Y_Actual, Y_Predicted, labels=np.unique(Y_Actual)))

# Using Seaborn heatmap to create the plot
fx = sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='d')

# labels the title and x, y axis of plot
fx.set_title('Plotting Confusion Matrix using Seaborn\n\n');
fx.set_xlabel('Predicted Values')
fx.set_ylabel('Actual Values ');

plt.show()
```

Plotting Confusion Matrix using Seaborn



In [183]:

```
from sklearn.metrics import classification_report
print(classification_report(Y_Actual, Y_Predicted))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	980
1	0.98	0.99	0.99	1135
2	0.98	0.97	0.97	1032
3	0.97	0.98	0.97	1010
4	0.98	0.96	0.97	982
5	0.97	0.97	0.97	892
6	0.98	0.97	0.97	958
7	0.97	0.96	0.97	1028
8	0.95	0.97	0.96	974
9	0.96	0.96	0.96	1009
accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000

In [184]:

```
unique_p, counts_p = np.unique(Y_Predicted, return_counts=True)
```

In [185]:

```
unique_p
```

Out[185]:

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

In [186]:

```
counts_p
```

Out[186]:

```
array([ 987, 1139, 1019, 1018,  966,  896,  949, 1021,  995, 1010])
```

REPORTING ACCURACY OF MODEL

TRAIN ACCURACY: 98.43%

TEST ACCURACY: 97.12%

Inference: This model performs better than baseline sigmoid and also better than tanh activation but looks like it is tending to overfit, regularization might help in this case.