

## Importing Required Libraries

In [78]:

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
from torch.utils.data.dataset import TensorDataset
from torch.utils.data import DataLoader
```

## DATASET LOADING AND PREPARATION

In [79]:

```
train_dataset = torchvision.datasets.MNIST(root='./data', train=True, transform=
transforms.ToTensor(), download=True)
test_dataset = torchvision.datasets.MNIST(root='./data', train=False, transform=
transforms.ToTensor())
```

In [80]:

```
loader_train = torch.utils.data.DataLoader(dataset=train_dataset, batch_size = le
n(train_dataset) ,shuffle=True)
loader_test = torch.utils.data.DataLoader(dataset=test_dataset, batch_size = len(
test_dataset) ,shuffle=True)
```

In [81]:

```
X_train,y_train=next(iter(loader_train))
X_test,y_test=next(iter(loader_test))
```

In [82]:

```
X_train=X_train.numpy()
y_train=y_train.numpy()
X_test=X_test.numpy()
y_test=y_test.numpy()
```

In [83]:

```
X_train_flattened=X_train.reshape(X_train.shape[0],X_train.shape[1]*X_train.shap
e[2]*X_train.shape[3])
X_test_flattened=X_test.reshape(X_test.shape[0],X_test.shape[1]*X_test.shape[2]*
X_test.shape[3])
```

## One-Hot Encoding

In [84]:

```
train_labels_encoded = []
for i in y_train:
    A=np.array([0]*10)
    A[i]=1
    train_labels_encoded.append(A)
y_train_encoded=np.array(train_labels_encoded)
```

In [85]:

```
test_labels_encoded = []
for i in y_test:
    A=np.array([0]*10)
    A[i]=1
    test_labels_encoded.append(A)
y_test_encoded=np.array(test_labels_encoded)
```

In [86]:

```
X_train_flattened_torch=torch.from_numpy(X_train_flattened)
y_train_encoded_torch=torch.from_numpy(y_train_encoded)
```

In [87]:

```
X_test_flattened_torch=torch.from_numpy(X_test_flattened)
y_test_encoded_torch=torch.from_numpy(y_test_encoded)
```

Here we have prepared test and train dataset with flattened images and one hot encoded labels

In [88]:

```
Train_Dataset=TensorDataset(X_train_flattened_torch,y_train_encoded_torch)
```

In [89]:

```
Test_Dataset=TensorDataset(X_test_flattened_torch,y_test_encoded_torch)
```

## HELPER FUNCTIONS

In [90]:

```
def relu(z):
    return np.maximum(0,z)

def softmax(z):
    return np.exp(z)/sum(np.exp(z))
```

In [91]:

```
def dif_relu(z):
    # return np.multiply(1.0 , (z>0))
    return z>0
```

In [92]:

```
def glorot_initialisation(output_n,input_n):
    M=np.sqrt(6/(input_n+output_n))
    W=np.random.uniform(low=-M, high=M, size=(output_n,input_n))
    b=np.random.uniform(low=-M, high=M, size=(output_n,1))
    return W,b
```

In [93]:

```
def initialize_parameters(layer_dims):
    """
    Arguments:
        layer_dims -- python array (list) containing the dimensions of each layer in
our network

    Returns:
        parameters -- python dictionary containing your parameters "W1", "b1", ...,
"WL", "bL":
            W1 -- weight matrix of shape (layer_dims[1], layer_dims[1-
1])
            b1 -- bias vector of shape (layer_dims[1], 1)
            Wl -- weight matrix of shape (layer_dims[l-1], layer_dims
[l])
            bl -- bias vector of shape (1, layer_dims[l])

    Tips:
        - For example: the layer_dims for the "Planar Data classification model" wou
ld have been [2,2,1].
        This means W1's shape was (2,2), b1 was (1,2), W2 was (2,1) and b2 was (1,
1). Now you have to generalize it!
        - In the for loop, use parameters['W' + str(l)] to access Wl, where l is the
iterative integer.
    """

    np.random.seed(1390)
    parameters = {}
    L = len(layer_dims) # number of layers in the network

    for l in range(1, L):
        # M=np.sqrt(6/(self.input_n+self.output_n))
        parameters['W' + str(l)],parameters['b' + str(l)] = glorot_initialisati
on(layer_dims[l], layer_dims[l-1])

        assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l-1
]))
        assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

    return parameters
```

In [94]:

```

def forward_propagation(X, parameters):
    """
    SHAPE OF X = 784,samples(i.e. 64 for a batch)
    Implements the forward propagation (and computes the loss) presented in Figure 2.

    Arguments:
    X -- input dataset, of shape (input size, number of examples)
    parameters -- python dictionary containing your parameters "W1", "b1", "W2", "b2", "W3", "b3", "W4", "b4":
        W1 -- weight matrix of shape (500,784)
        b1 -- bias vector of shape (500,1)
        W2 -- weight matrix of shape (250,500)
        b2 -- bias vector of shape (250,1)
        W3 -- weight matrix of shape (100,250)
        b3 -- bias vector of shape (100,1)
        W4 -- weight matrix of shape (10,100)
        b4 -- bias vector of shape (10,1)

    Returns:
    loss -- the loss function (vanilla logistic loss)
    """

    cache={}

    cache['Z2']=np.dot(parameters['W1'],X)+parameters['b1']
    cache['A2']=relu(cache['Z2'])

    cache['Z3']=np.dot(parameters['W2'],cache['A2'])+parameters['b2']
    cache['A3']=relu(cache['Z3'])

    cache['Z4']=np.dot(parameters['W3'],cache['A3'])+parameters['b3']
    cache['A4']=relu(cache['Z4'])

    cache['Z5']=np.dot(parameters['W4'],cache['A4'])+parameters['b4']
    cache['A5']=softmax(cache['Z5'])

    A_last=cache['A5']

    cache['W1']=parameters['W1']
    cache['b1']=parameters['b1']

    cache['W2']=parameters['W2']
    cache['b2']=parameters['b2']

    cache['W3']=parameters['W3']
    cache['b3']=parameters['b3']

    cache['W4']=parameters['W4']
    cache['b4']=parameters['b4']

    # cache = {"W1": W1, "b1": b1, "Z2": Z2, "A2": A2,
    #           "W2": W2, "b2": b2, "Z3": Z3, "A3": A3,
    #           "W3": W3, "b3": b3, "Z4": Z4, "A4": A4,
    #           "W4": W4, "b4": b4, "Z5": Z5, "A5": A5}

```

```
return A_last, cache
```

In [95]:

```
# Calculating the loss function using the cross entropy
"""Arguments:
    A -- post-activation, output of forward propagation
    Y -- "true" labels vector, same shape as A

Returns:
    cost - value of the cost function
"""
def compute_cost(A, Y, cache, lambd):
    #A is predicted
    #Y is actual
    m = Y.shape[1]
    logprobs = np.multiply(-np.log(A), Y) + np.multiply(-np.log(1 - A), 1 - Y)
    cost = 1./m * np.nansum(logprobs)

    l2_regularisation_cost=(lambd/(2*m))*(np.sum(np.square(cache['W1']))+np.sum(
np.square(cache['W2']))+np.sum(np.square(cache['W3']))+np.sum(np.square(cache['W
4'])))

    regularized_total_cost=cost +l2_regularisation_cost
    return regularized_total_cost
```

In [96]:

```
def backward_propagation(X, Y, cache, lambd):

    m=X.shape[1]

    grads={}

    grads['dz5']=cache['A5']-Y
    grads['dw4']= 1./m * np.dot(grads['dz5'],cache['A4'].T) + (lambd*cache['W4'])/
m
    grads['db4']= 1./m * np.sum(grads['dz5'],axis=1,keepdims=True)

    grads['dA4']=np.dot(cache['W4'].T,grads['dz5'])
    grads['dz4']=np.multiply(grads['dA4'],dif_relu(cache['Z4']))
    grads['dw3']=1./m * np.dot(grads['dz4'],cache['A3'].T) + (lambd*cache['W3'])/m
    grads['db3']=1./m * np.sum(grads['dz4'],axis=1,keepdims=True)

    grads['dA3']=np.dot(cache['W3'].T,grads['dz4'])
    grads['dz3']=np.multiply(grads['dA3'],dif_relu(cache['Z3']))
    grads['dw2']=1./m * np.dot(grads['dz3'],cache['A2'].T) + (lambd*cache['W2'])/
m
    grads['db2']=1./m * np.sum(grads['dz3'],axis=1,keepdims=True)

    grads['dA2']=np.dot(cache['W2'].T,grads['dz3'])
    grads['dz2']=np.multiply(grads['dA2'],dif_relu(cache['Z2']))
    grads['dw1']=1./m * np.dot(grads['dz2'],X.T) + (lambd*cache['W1'])/m
    grads['db1']=1./m * np.sum(grads['dz2'],axis=1,keepdims=True)

    return grads
```

In [97]:

```
def update_parameters(parameters, grads, learning_rate):

    updated_parameters={}

    updated_parameters['W1']=parameters['W1']-learning_rate*grads['dW1']
    updated_parameters['b1']=parameters['b1']-learning_rate*grads['db1']

    updated_parameters['W2']=parameters['W2']-learning_rate*grads['dW2']
    updated_parameters['b2']=parameters['b2']-learning_rate*grads['db2']

    updated_parameters['W3']=parameters['W3']-learning_rate*grads['dW3']
    updated_parameters['b3']=parameters['b3']-learning_rate*grads['db3']

    updated_parameters['W4']=parameters['W4']-learning_rate*grads['dW4']
    updated_parameters['b4']=parameters['b4']-learning_rate*grads['db4']

    return updated_parameters
```

In [98]:

```
#Finding the accuracy of the parameter at the output
"""
Arguments:
y_actual - given in the dataset / also called as the ground truth
y_pred - generated from the neural network , after a series of forward and b
ackprop

Returns:
accuracy = finding the matches of the predicted vs the actual
"""
def calculate_accuracy(y_actual,y_pred):
    accuracy = np.count_nonzero(np.argmax(y_pred,axis=0)==np.argmax(y_actual,axis=0))/y_actual.shape[1]
    return accuracy
```

In [99]:

```
def predict(X,Y,parameters):

    """
This function is used to predict the results of a n-layer neural network.

Arguments:
X -- data set of examples you would like to label
Y -- data set of examples
parameters -- parameters of the trained model

Returns:
ypred -- predictions for the given dataset X
"""

    y_pred,cache=forward_propagation(X,parameters)
    return y_pred
```

## MODEL TRAINING

In [100]:

```
def model(Train_Dataset,layer_dimensions,total_epochs=15,Batch_Size=64,learning_
rate=0.01,lambd=0.8):

    costs=[]
    accuracy=[]

    parameters=initialize_parameters(layer_dimensions)
    num_iterations=len(Train_Dataset)//Batch_Size

    #Train_Dataset=TensorDataset(X_training,Y_training)
    for epoch in range(total_epochs):
        for iteration in range(num_iterations):
            Data_Loader=torch.utils.data.DataLoader(dataset=Train_Dataset,batch_size=6
4, shuffle=True)

            data_iter=iter(Data_Loader)
            Data=next(data_iter)
            X,y=Data #X.shape=(batch_size,784) y.shape=(batch_size,10)
            X=X.numpy()
            y=y.numpy()
            a5,cache=forward_propagation(X.T,parameters)
            cost=compute_cost(a5,y.T,cache,lambd)
            gradients=backward_propagation(X.T,y.T,cache,lambd)
            parameters=update_parameters(parameters,gradients,learning_rate)

            if iteration%200==0:
                print("epoch: ",epoch+1,"/",total_epochs, "  iteration= ",iteration+1,
"/",num_iterations, "  Loss: ",cost)
                accuracy.append(calculate_accuracy(y.T,a5))
                costs.append(cost)

    return accuracy, costs, parameters
```

In [114]:

```
layer_dimensions=[784,500,250,100,10]  
Train_accuracy,Train_costs,Trained_parameters=model(Train_Dataset,layer_dimensions,15,64,0.01,0.8)
```



```
epoch: 1 / 15 iteration= 1 / 937 Loss: 10.24665749963704
epoch: 1 / 15 iteration= 201 / 937 Loss: 8.564261611610512
epoch: 1 / 15 iteration= 401 / 937 Loss: 7.420605826921399
epoch: 1 / 15 iteration= 601 / 937 Loss: 6.885346259516561
epoch: 1 / 15 iteration= 801 / 937 Loss: 6.403292172120739
epoch: 2 / 15 iteration= 1 / 937 Loss: 6.452327503465423
epoch: 2 / 15 iteration= 201 / 937 Loss: 6.089522975943435
epoch: 2 / 15 iteration= 401 / 937 Loss: 5.962707372903211
epoch: 2 / 15 iteration= 601 / 937 Loss: 5.4670342284631905
epoch: 2 / 15 iteration= 801 / 937 Loss: 5.202636592920128
epoch: 3 / 15 iteration= 1 / 937 Loss: 4.77820843767748
epoch: 3 / 15 iteration= 201 / 937 Loss: 4.841851990559459
epoch: 3 / 15 iteration= 401 / 937 Loss: 4.696018944101629
epoch: 3 / 15 iteration= 601 / 937 Loss: 4.488708957460135
epoch: 3 / 15 iteration= 801 / 937 Loss: 4.28934694401448
epoch: 4 / 15 iteration= 1 / 937 Loss: 4.028385348035927
epoch: 4 / 15 iteration= 201 / 937 Loss: 4.0758647084738895
epoch: 4 / 15 iteration= 401 / 937 Loss: 3.974530131459747
epoch: 4 / 15 iteration= 601 / 937 Loss: 3.6465236357136117
epoch: 4 / 15 iteration= 801 / 937 Loss: 3.637218839290589
epoch: 5 / 15 iteration= 1 / 937 Loss: 3.734181223285472
epoch: 5 / 15 iteration= 201 / 937 Loss: 3.5855167756247734
epoch: 5 / 15 iteration= 401 / 937 Loss: 3.2157957371924164
epoch: 5 / 15 iteration= 601 / 937 Loss: 3.0251044494455823
epoch: 5 / 15 iteration= 801 / 937 Loss: 3.0387374471144772
epoch: 6 / 15 iteration= 1 / 937 Loss: 2.92617117637501
epoch: 6 / 15 iteration= 201 / 937 Loss: 2.851057199354434
epoch: 6 / 15 iteration= 401 / 937 Loss: 2.7548506878929024
epoch: 6 / 15 iteration= 601 / 937 Loss: 2.6273646131420714
epoch: 6 / 15 iteration= 801 / 937 Loss: 2.655667992753477
epoch: 7 / 15 iteration= 1 / 937 Loss: 2.51058362023355
epoch: 7 / 15 iteration= 201 / 937 Loss: 2.3837788372842788
epoch: 7 / 15 iteration= 401 / 937 Loss: 2.3898939307404827
epoch: 7 / 15 iteration= 601 / 937 Loss: 2.243522821889658
epoch: 7 / 15 iteration= 801 / 937 Loss: 2.3308796751895793
epoch: 8 / 15 iteration= 1 / 937 Loss: 2.1649825337329487
epoch: 8 / 15 iteration= 201 / 937 Loss: 2.0946083964101474
epoch: 8 / 15 iteration= 401 / 937 Loss: 2.2828127545878028
epoch: 8 / 15 iteration= 601 / 937 Loss: 1.8249833646968407
epoch: 8 / 15 iteration= 801 / 937 Loss: 1.7898704587402183
epoch: 9 / 15 iteration= 1 / 937 Loss: 1.8255597629277764
epoch: 9 / 15 iteration= 201 / 937 Loss: 1.9323719681470295
epoch: 9 / 15 iteration= 401 / 937 Loss: 1.7352646031430177
epoch: 9 / 15 iteration= 601 / 937 Loss: 1.8296360457139014
epoch: 9 / 15 iteration= 801 / 937 Loss: 1.6330053947461676
epoch: 10 / 15 iteration= 1 / 937 Loss: 1.6274586381869076
epoch: 10 / 15 iteration= 201 / 937 Loss: 1.6261009944552596
epoch: 10 / 15 iteration= 401 / 937 Loss: 1.4480423532469877
epoch: 10 / 15 iteration= 601 / 937 Loss: 1.5667137942252174
epoch: 10 / 15 iteration= 801 / 937 Loss: 1.444136211820898
epoch: 11 / 15 iteration= 1 / 937 Loss: 1.3858226251440717
epoch: 11 / 15 iteration= 201 / 937 Loss: 1.6059569652673968
epoch: 11 / 15 iteration= 401 / 937 Loss: 1.3160241285108274
epoch: 11 / 15 iteration= 601 / 937 Loss: 1.2525268225964128
epoch: 11 / 15 iteration= 801 / 937 Loss: 1.3333353130193764
epoch: 12 / 15 iteration= 1 / 937 Loss: 1.3257485188478042
epoch: 12 / 15 iteration= 201 / 937 Loss: 1.209194326008392
epoch: 12 / 15 iteration= 401 / 937 Loss: 1.1110561832529737
epoch: 12 / 15 iteration= 601 / 937 Loss: 1.238720947743997
epoch: 12 / 15 iteration= 801 / 937 Loss: 1.2932704654218936
epoch: 13 / 15 iteration= 1 / 937 Loss: 1.464326811340972
```


```
epoch: 13 / 15 iteration= 201 / 937 Loss: 1.272144386531624
epoch: 13 / 15 iteration= 401 / 937 Loss: 1.3448446867930093
epoch: 13 / 15 iteration= 601 / 937 Loss: 1.1313718587661445
epoch: 13 / 15 iteration= 801 / 937 Loss: 1.2057377839806718
epoch: 14 / 15 iteration= 1 / 937 Loss: 1.0369554623507273
epoch: 14 / 15 iteration= 201 / 937 Loss: 1.0573854059126497
epoch: 14 / 15 iteration= 401 / 937 Loss: 1.2270571120069174
epoch: 14 / 15 iteration= 601 / 937 Loss: 0.8787962443751072
epoch: 14 / 15 iteration= 801 / 937 Loss: 1.1267776343406106
epoch: 15 / 15 iteration= 1 / 937 Loss: 1.1392730653798524
epoch: 15 / 15 iteration= 201 / 937 Loss: 1.1172321624453316
epoch: 15 / 15 iteration= 401 / 937 Loss: 0.8915757681278184
epoch: 15 / 15 iteration= 601 / 937 Loss: 0.9160363899972161
epoch: 15 / 15 iteration= 801 / 937 Loss: 0.9858567924604414
```

In [102]:

```
plt.plot(Train_costs)
```

Out[102]:

[<matplotlib.lines.Line2D at 0x7fb97f6a7810>]




In [103]:

```
plt.plot(Train_accuracy)
```

Out[103]:

[<matplotlib.lines.Line2D at 0x7fb97f687950>]



In [104]:

```
Train_accuracy[-1]
```

Out[104]:

0.921875

In [105]:

```
Test_Data_Loader=torch.utils.data.DataLoader(dataset=Test_Dataset,batch_size=len
(Test_Dataset),shuffle=True)
data_iter=iter(Test_Data_Loader)
Test_Data=next(data_iter)
X,y=Test_Data
X=X.numpy()
y=y.numpy()
y_predicted,cache_out=forward_propagation(X.T,Trained_parameters)
Test_accuracy=calculate_accuracy(y.T,y_predicted)
```

In [106]:

```
Test_accuracy
```

Out[106]:

0.9368

In [107]:

```
Y_Predicted=np.array(np.argmax(y_predicted,axis=0))
```

In [108]:

```
Y_Actual=np.array(np.argmax(y.T,axis=0) )
```

## CONFUSION MATRIX AND CLASSIFICATION REPORT

In [109]:


```
from sklearn.metrics import confusion_matrix
import seaborn as sns
plt.figure(figsize=(10,10))
conf_matrix = (confusion_matrix(Y_Actual, Y_Predicted, labels=np.unique(Y_Actual)))

# Using Seaborn heatmap to create the plot
fx = sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='d')

# labels the title and x, y axis of plot
fx.set_title('Plotting Confusion Matrix using Seaborn\n\n');
fx.set_xlabel('Predicted Values')
fx.set_ylabel('Actual Values ');

plt.show()
```

Plotting Confusion Matrix using Seaborn



In [110]:

```
from sklearn.metrics import classification_report
print(classification_report(Y_Actual, Y_Predicted))
```

	precision	recall	f1-score	support
0	0.95	0.98	0.96	980
1	0.97	0.98	0.97	1135
2	0.94	0.92	0.93	1032
3	0.92	0.94	0.93	1010
4	0.90	0.96	0.93	982
5	0.95	0.90	0.92	892
6	0.92	0.95	0.94	958
7	0.96	0.92	0.94	1028
8	0.93	0.91	0.92	974
9	0.93	0.90	0.92	1009
accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000

In [111]:

```
unique_p, counts_p = np.unique(Y_Predicted, return_counts=True)
```

In [112]:

```
unique_p
```

Out[112]:

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

In [113]:

```
counts_p
```

Out[113]:

```
array([1009, 1152, 1016, 1030, 1040, 842, 991, 988, 949, 983])
```

## REPORTING ACCURACY OF MODEL

TRAIN ACCURACY: 92.18%

TEST ACCURACY: 93.68%

lambda=0.8

*Inference: This model performs outstands all unregularised models with sigmoid, tanh, and relu activations. This shows that regularisation can help in model improvement.*