

Importing Required Libraries

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
from torch.utils.data.dataset import TensorDataset
from torch.utils.data import DataLoader
```

DATASET LOADING AND PREPARATION

In [2]:

```
train_dataset = torchvision.datasets.MNIST(root='./data', train=True, transform=
transforms.ToTensor(), download=True)
test_dataset = torchvision.datasets.MNIST(root='./data', train=False, transform=
transforms.ToTensor())
```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

In [3]:

```
loader_train = torch.utils.data.DataLoader(dataset=train_dataset, batch_size = len(train_dataset), shuffle=True)
loader_test = torch.utils.data.DataLoader(dataset=test_dataset, batch_size = len(test_dataset), shuffle=True)
```

In [4]:

```
X_train, y_train = next(iter(loader_train))
X_test, y_test = next(iter(loader_test))
```

In [5]:

```
X_train = X_train.numpy()
y_train = y_train.numpy()
X_test = X_test.numpy()
y_test = y_test.numpy()
```

In [6]:

```
X_train_flattened = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2]*X_train.shape[3])
X_test_flattened = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2]*X_test.shape[3])
```

One-Hot Encoding

In [7]:

```
train_labels_encoded = []
for i in y_train:
    A = np.array([0]*10)
    A[i] = 1
    train_labels_encoded.append(A)
y_train_encoded = np.array(train_labels_encoded)
```

In [8]:

```
test_labels_encoded = []
for i in y_test:
    A = np.array([0]*10)
    A[i] = 1
    test_labels_encoded.append(A)
y_test_encoded = np.array(test_labels_encoded)
```

In [9]:

```
X_train_flattened_torch = torch.from_numpy(X_train_flattened)
y_train_encoded_torch = torch.from_numpy(y_train_encoded)
```

In [10]:

```
X_test_flattened_torch = torch.from_numpy(X_test_flattened)
y_test_encoded_torch = torch.from_numpy(y_test_encoded)
```

Here we have prepared dataset with flattened images and one-hot encoded labels

In [11]:

```
Train_Dataset=TensorDataset(X_train_flattened_torch,y_train_encoded_torch)
```

In [12]:

```
Test_Dataset=TensorDataset(X_test_flattened_torch,y_test_encoded_torch)
```

HELPER FUNCTIONS

In [13]:

```
def sigmoid(z):  
    return 1/(1+(np.exp(-z)))  
  
def softmax(z):  
    return np.exp(z)/sum(np.exp(z))
```

In [14]:

```
def dif_sigmoid(z):  
    return ((1-sigmoid(z))*sigmoid(z)) ##since diff of sigmoid is (1-sigmoid)*sigmoid
```

In [15]:

```
def glorot_initialisation(output_n,input_n):  
    M=np.sqrt(6/(input_n+output_n))  
    W=np.random.uniform(low=-M, high=M, size=(output_n,input_n))  
    b=np.random.uniform(low=-M, high=M, size=(output_n,1))  
    return W,b
```

In [16]:

```

def initialize_parameters(layer_dims):
    """
    Arguments:
        layer_dims -- python array (list) containing the dimensions of each layer in
our network

    Returns:
        parameters -- python dictionary containing your parameters "W1", "b1", ...,
"WL", "bL":
            W1 -- weight matrix of shape (layer_dims[l], layer_dims[l-
1])
            b1 -- bias vector of shape (layer_dims[l], 1)
            Wl -- weight matrix of shape (layer_dims[l-1], layer_dims
[l])
            bl -- bias vector of shape (1, layer_dims[l])

    Tips:
        - For example: the layer_dims for the "Planar Data classification model" wou
ld have been [2,2,1].
        This means W1's shape was (2,2), b1 was (1,2), W2 was (2,1) and b2 was (1,
1). Now you have to generalize it!
        - In the for loop, use parameters['W' + str(l)] to access Wl, where l is the
iterative integer.
    """

    np.random.seed(1390)
    parameters = {}
    L = len(layer_dims) # number of layers in the network

    for l in range(1, L):

        # M=np.sqrt(6/(self.input_n+self.output_n))
        parameters['W' + str(l)], parameters['b' + str(l)] = glorot_initialisati
on(layer_dims[l], layer_dims[l-1])

        assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l-1
]))
        assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

    return parameters

```

In [17]:

```

def forward_propagation(X, parameters):
    """
    SHAPE OF X = 784,samples(i.e. 64 for a batch)
    Implements the forward propagation (and computes the loss) presented in Figure 2.

    Arguments:
    X -- input dataset, of shape (input size, number of examples)
    parameters -- python dictionary containing your parameters "W1", "b1", "W2", "b2", "W3", "b3", "W4", "b4":
        W1 -- weight matrix of shape (500,784)
        b1 -- bias vector of shape (500,1)
        W2 -- weight matrix of shape (250,500)
        b2 -- bias vector of shape (250,1)
        W3 -- weight matrix of shape (100,250)
        b3 -- bias vector of shape (100,1)
        W4 -- weight matrix of shape (10,100)
        b4 -- bias vector of shape (10,1)

    Returns:
    loss -- the loss function (vanilla logistic loss)
    """

    cache={}

    cache['Z2']=np.dot(parameters['W1'],X)+parameters['b1']
    cache['A2']=sigmoid(cache['Z2'])

    cache['Z3']=np.dot(parameters['W2'],cache['A2'])+parameters['b2']
    cache['A3']=sigmoid(cache['Z3'])

    cache['Z4']=np.dot(parameters['W3'],cache['A3'])+parameters['b3']
    cache['A4']=sigmoid(cache['Z4'])

    cache['Z5']=np.dot(parameters['W4'],cache['A4'])+parameters['b4']
    cache['A5']=softmax(cache['Z5'])

    A_last=cache['A5']

    cache['W1']=parameters['W1']
    cache['b1']=parameters['b1']

    cache['W2']=parameters['W2']
    cache['b2']=parameters['b2']

    cache['W3']=parameters['W3']
    cache['b3']=parameters['b3']

    cache['W4']=parameters['W4']
    cache['b4']=parameters['b4']

    # cache = {"W1": W1, "b1": b1, "Z2": Z2, "A2": A2,
    #           "W2": W2, "b2": b2, "Z3": Z3, "A3": A3,
    #           "W3": W3, "b3": b3, "Z4": Z4, "A4": A4,
    #           "W4": W4, "b4": b4, "Z5": Z5, "A5": A5}

```

```
return A_last, cache
```

In [18]:

```
# Calculating the loss function using the cross entropy
"""Arguments:
    A -- post-activation, output of forward propagation
    Y -- "true" labels vector, same shape as A

Returns:
    cost - value of the cost function
"""
def compute_cost(A, Y):
    #A is predicted
    #Y is actual
    m = Y.shape[1]
    logprobs = np.multiply(-np.log(A), Y) + np.multiply(-np.log(1 - A), 1 - Y)
    cost = 1./m * np.nansum(logprobs)
    return cost
```

In [19]:

```
def backward_propagation(X, Y, cache):

    m=X.shape[1]

    grads={}

    grads['dz5']=cache['A5']-Y
    grads['dw4']= 1./m * np.dot(grads['dz5'],cache['A4'].T)
    grads['db4']= 1./m * np.sum(grads['dz5'],axis=1,keepdims=True)

    grads['dA4']=np.dot(cache['W4'].T,grads['dz5'])
    grads['dz4']=np.multiply(grads['dA4'],dif_sigmoid(cache['z4']))
    grads['dw3']=1./m * np.dot(grads['dz4'],cache['A3'].T)
    grads['db3']=1./m * np.sum(grads['dz4'],axis=1,keepdims=True)

    grads['dA3']=np.dot(cache['W3'].T,grads['dz4'])
    grads['dz3']=np.multiply(grads['dA3'],dif_sigmoid(cache['z3']))
    grads['dw2']=1./m * np.dot(grads['dz3'],cache['A2'].T)
    grads['db2']=1./m * np.sum(grads['dz3'],axis=1,keepdims=True)

    grads['dA2']=np.dot(cache['W2'].T,grads['dz3'])
    grads['dz2']=np.multiply(grads['dA2'],dif_sigmoid(cache['z2']))
    grads['dw1']=1./m * np.dot(grads['dz2'],X.T)
    grads['db1']=1./m * np.sum(grads['dz2'],axis=1,keepdims=True)

    return grads
```

In [20]:

```
def update_parameters(parameters, grads, learning_rate):

    updated_parameters={}

    updated_parameters['W1']=parameters['W1']-learning_rate*grads['dW1']
    updated_parameters['b1']=parameters['b1']-learning_rate*grads['db1']

    updated_parameters['W2']=parameters['W2']-learning_rate*grads['dW2']
    updated_parameters['b2']=parameters['b2']-learning_rate*grads['db2']

    updated_parameters['W3']=parameters['W3']-learning_rate*grads['dW3']
    updated_parameters['b3']=parameters['b3']-learning_rate*grads['db3']

    updated_parameters['W4']=parameters['W4']-learning_rate*grads['dW4']
    updated_parameters['b4']=parameters['b4']-learning_rate*grads['db4']

    return updated_parameters
```

In [21]:

```
#Finding the accuracy of the parameter at the output
"""
    Arguments:
    y_actual - given in the dataset / also called as the ground truth
    y_pred - generated from the neural network , after a series of forward and b
ackprop

    Returns:
    accuracy = finding the matches of the prdicted vs the actual
"""
def calculate_accuracy(y_actual,y_pred):
    accuracy = np.count_nonzero(np.argmax(y_pred,axis=0)==np.argmax(y_actual,axis=0))/y_actual.shape[1]
    return accuracy
```

In [22]:

```
def predict(X,Y,parameters):

    """
    This function is used to predict the results of a n-layer neural network.

    Arguments:
    X -- data set of examples you would like to label
    Y -- data set of examples
    parameters -- parameters of the trained model

    Returns:
    ypred -- predictions for the given dataset X
    """

    y_pred,cache=forward_propagation(X,parameters)
    return y_pred
```

MODEL TRAINING

In [23]:

```
def model(Train_Dataset,layer_dimensions,total_epochs=15,Batch_Size=64,learning_
rate=0.01):

    costs=[]
    accuracy=[]

    parameters=initialize_parameters(layer_dimensions)
    num_iterations=len(Train_Dataset)//Batch_Size

    for epoch in range(total_epochs):
        for iteration in range(num_iterations):
            Data_Loader=torch.utils.data.DataLoader(dataset=Train_Dataset,batch_size=6
4, shuffle=True)

            data_iter=iter(Data_Loader)
            Data=next(data_iter)
            X,y=Data #X.shape=(batch_size,784) y.shape=(batch_size,10)
            X=X.numpy()
            y=y.numpy()
            a5,cache=forward_propagation(X.T,parameters)
            cost=compute_cost(a5,y.T)
            gradients=backward_propagation(X.T,y.T,cache)
            parameters=update_parameters(parameters,gradients,learning_rate)

            if iteration%200==0:
                print("epoch: ",epoch+1,"/",total_epochs, " iteration= ",iteration+1,
"/",num_iterations, " Loss: ",cost)
                accuracy.append(calculate_accuracy(y.T,a5))
                costs.append(cost)

    return accuracy,costs,parameters
```


In [24]:

```
layer_dimensions=[784,500,250,100,10]  
Train_accuracy,Train_costs,Trained_parameters=model(Train_Dataset,layer_dimensions,15,64,0.01)
```

```
epoch: 1 / 15 iteration= 1 / 937 Loss: 3.528933913473381
epoch: 1 / 15 iteration= 201 / 937 Loss: 3.2295368549473595
epoch: 1 / 15 iteration= 401 / 937 Loss: 3.242631231086987
epoch: 1 / 15 iteration= 601 / 937 Loss: 3.2512596589972222
epoch: 1 / 15 iteration= 801 / 937 Loss: 3.2301902686212207
epoch: 2 / 15 iteration= 1 / 937 Loss: 3.2447326250897324
epoch: 2 / 15 iteration= 201 / 937 Loss: 3.2324886265456714
epoch: 2 / 15 iteration= 401 / 937 Loss: 3.2333129665809
epoch: 2 / 15 iteration= 601 / 937 Loss: 3.240644758292385
epoch: 2 / 15 iteration= 801 / 937 Loss: 3.2188152128410774
epoch: 3 / 15 iteration= 1 / 937 Loss: 3.202937459221192
epoch: 3 / 15 iteration= 201 / 937 Loss: 3.2384451696367775
epoch: 3 / 15 iteration= 401 / 937 Loss: 3.2243901901273264
epoch: 3 / 15 iteration= 601 / 937 Loss: 3.21005947468937
epoch: 3 / 15 iteration= 801 / 937 Loss: 3.196689693254491
epoch: 4 / 15 iteration= 1 / 937 Loss: 3.203091746347312
epoch: 4 / 15 iteration= 201 / 937 Loss: 3.2120254483484327
epoch: 4 / 15 iteration= 401 / 937 Loss: 3.1849661365464756
epoch: 4 / 15 iteration= 601 / 937 Loss: 3.194624929340704
epoch: 4 / 15 iteration= 801 / 937 Loss: 3.179794590903228
epoch: 5 / 15 iteration= 1 / 937 Loss: 3.154745830565319
epoch: 5 / 15 iteration= 201 / 937 Loss: 3.112783349933035
epoch: 5 / 15 iteration= 401 / 937 Loss: 3.1595545330993455
epoch: 5 / 15 iteration= 601 / 937 Loss: 3.152586387242789
epoch: 5 / 15 iteration= 801 / 937 Loss: 3.1192595486985772
epoch: 6 / 15 iteration= 1 / 937 Loss: 3.1017414402166676
epoch: 6 / 15 iteration= 201 / 937 Loss: 3.1110908336395413
epoch: 6 / 15 iteration= 401 / 937 Loss: 3.059002938657069
epoch: 6 / 15 iteration= 601 / 937 Loss: 3.0556463153629165
epoch: 6 / 15 iteration= 801 / 937 Loss: 2.996314992931155
epoch: 7 / 15 iteration= 1 / 937 Loss: 2.9535544743377105
epoch: 7 / 15 iteration= 201 / 937 Loss: 2.893522977960649
epoch: 7 / 15 iteration= 401 / 937 Loss: 2.889575061301709
epoch: 7 / 15 iteration= 601 / 937 Loss: 2.7733852872204485
epoch: 7 / 15 iteration= 801 / 937 Loss: 2.7178981347073625
epoch: 8 / 15 iteration= 1 / 937 Loss: 2.6631192335457547
epoch: 8 / 15 iteration= 201 / 937 Loss: 2.5806944120456015
epoch: 8 / 15 iteration= 401 / 937 Loss: 2.5204179631551566
epoch: 8 / 15 iteration= 601 / 937 Loss: 2.4621674905573956
epoch: 8 / 15 iteration= 801 / 937 Loss: 2.4706587507795685
epoch: 9 / 15 iteration= 1 / 937 Loss: 2.1807854558629143
epoch: 9 / 15 iteration= 201 / 937 Loss: 2.3088419527638417
epoch: 9 / 15 iteration= 401 / 937 Loss: 2.2206123741188692
epoch: 9 / 15 iteration= 601 / 937 Loss: 2.0933125807780755
epoch: 9 / 15 iteration= 801 / 937 Loss: 2.0711535664981566
epoch: 10 / 15 iteration= 1 / 937 Loss: 2.1878524693648878
epoch: 10 / 15 iteration= 201 / 937 Loss: 1.9416069624938286
epoch: 10 / 15 iteration= 401 / 937 Loss: 1.96478963780603
epoch: 10 / 15 iteration= 601 / 937 Loss: 1.8452532111874778
epoch: 10 / 15 iteration= 801 / 937 Loss: 1.8312166271346975
epoch: 11 / 15 iteration= 1 / 937 Loss: 1.658432782796748
epoch: 11 / 15 iteration= 201 / 937 Loss: 1.7922312269088874
epoch: 11 / 15 iteration= 401 / 937 Loss: 1.5365322614449102
epoch: 11 / 15 iteration= 601 / 937 Loss: 1.5530971284019228
epoch: 11 / 15 iteration= 801 / 937 Loss: 1.45867301401834
epoch: 12 / 15 iteration= 1 / 937 Loss: 1.6461964709829917
epoch: 12 / 15 iteration= 201 / 937 Loss: 1.7546243695318067
epoch: 12 / 15 iteration= 401 / 937 Loss: 1.508935180153573
epoch: 12 / 15 iteration= 601 / 937 Loss: 1.5657814183237377
epoch: 12 / 15 iteration= 801 / 937 Loss: 1.2779010605014365
epoch: 13 / 15 iteration= 1 / 937 Loss: 1.3677715753119157
```


```
epoch: 13 / 15    iteration= 201 / 937    Loss: 1.3593053619671391
epoch: 13 / 15    iteration= 401 / 937    Loss: 1.5256851323295058
epoch: 13 / 15    iteration= 601 / 937    Loss: 1.1525399712001616
epoch: 13 / 15    iteration= 801 / 937    Loss: 1.2509171601790114
epoch: 14 / 15    iteration= 1 / 937     Loss: 1.307047368715477
epoch: 14 / 15    iteration= 201 / 937    Loss: 1.2819540625030836
epoch: 14 / 15    iteration= 401 / 937    Loss: 1.0567583539852194
epoch: 14 / 15    iteration= 601 / 937    Loss: 1.2239266930802515
epoch: 14 / 15    iteration= 801 / 937    Loss: 1.1545694960790631
epoch: 15 / 15    iteration= 1 / 937     Loss: 1.120516714218322
epoch: 15 / 15    iteration= 201 / 937    Loss: 1.058407912053907
epoch: 15 / 15    iteration= 401 / 937    Loss: 1.0835894625660238
epoch: 15 / 15    iteration= 601 / 937    Loss: 1.2591326387087962
epoch: 15 / 15    iteration= 801 / 937    Loss: 1.3880216242889456
```

In [25]:

```
plt.plot(Train_costs)
```

Out[25]:

[<matplotlib.lines.Line2D at 0x7fc45f7b8ed0>]




In [26]:

```
plt.plot(Train_accuracy)
```

Out[26]:

[<matplotlib.lines.Line2D at 0x7fc45f6ddc50>]



In [27]:

```
Train_accuracy[-1]
```

Out[27]:

0.875

In [28]:

```
Test_Data_Loader=torch.utils.data.DataLoader(dataset=Test_Dataset,batch_size=len
(Test_Dataset),shuffle=True)
data_iter=iter(Test_Data_Loader)
Test_Data=next(data_iter)
X,y=Test_Data #X.shape=(batch_size,784) y.shape=(batch_size,10)
X=X.numpy()
y=y.numpy()
y_predicted,cache_out=forward_propagation(X.T,Trained_parameters)
cost=compute_cost(y_predicted,y.T)
Test_accuracy=calculate_accuracy(y.T,y_predicted)
```

In [29]:

```
Test_accuracy
```

Out[29]:

0.8223

In [30]:

```
Y_Predicted=np.array(np.argmax(y_predicted,axis=0))
```

In [31]:

```
Y_Actual=np.array(np.argmax(y.T,axis=0) )
```

CONFUSION MATRIX AND CLASSIFICATION REPORT

In [32]:

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
plt.figure(figsize=(10,10))
conf_matrix = (confusion_matrix(Y_Actual, Y_Predicted, labels=np.unique(Y_Actual)))


# Using Seaborn heatmap to create the plot
fx = sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='d')

# labels the title and x, y axis of plot
fx.set_title('Plotting Confusion Matrix using Seaborn\n\n');
fx.set_xlabel('Predicted Values')
fx.set_ylabel('Actual Values ');

plt.show()

```

Plotting Confusion Matrix using Seaborn



In [33]:

```
from sklearn.metrics import classification_report
print(classification_report(Y_Actual, Y_Predicted))
```

	precision	recall	f1-score	support
0	0.91	0.95	0.93	980
1	0.89	0.98	0.93	1135
2	0.86	0.81	0.83	1032
3	0.76	0.85	0.80	1010
4	0.74	0.79	0.77	982
5	0.75	0.62	0.68	892
6	0.85	0.89	0.87	958
7	0.86	0.89	0.87	1028
8	0.81	0.71	0.76	974
9	0.75	0.68	0.72	1009
accuracy				0.82
macro avg				0.82
weighted avg				0.82

In [34]:

```
unique_p, counts_p = np.unique(Y_Predicted, return_counts=True)
```

In [35]:

```
unique_p
```

Out[35]:

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

In [36]:

```
counts_p
```

Out[36]:

```
array([1031, 1244, 979, 1122, 1052, 738, 1014, 1055, 852, 913])
```

REPORTING ACCURACY OF MODEL

TRAIN ACCURACY: 87.5%

TEST ACCURACY: 82.23%